



# 10-701 Introduction to Machine Learning

---

## Deep Learning

### Readings:

Bishop Ch. 4.1.7, Ch. 5

Murphy Ch. 16.5, Ch. 28

Mitchell Ch. 4

Matt Gormley  
Lecture 13  
October 19, 2016

# Reminders

- Homework 3:
  - due 10/24/16

# Outline

- **Deep Neural Networks (DNNs)**
  - Three ideas for training a DNN
  - Experiments: MNIST digit classification
  - Autoencoders
  - Pretraining
- **Recurrent Neural Networks (RNNs)**
  - Bidirectional RNNs
  - Deep Bidirectional RNNs
  - Deep Bidirectional LSTMs
  - Connection to forward-backward algorithm
- **Convolutional Neural Networks (CNNs)**
  - Convolutional layers
  - Pooling layers
  - Image recognition

# **PRE-TRAINING FOR DEEP NETS**

## Goals for Today's Lecture

1.
  1. Explore a **new class of decision functions** (Deep Neural Networks)
  2. Consider **variants of this recipe** for training

2. Choose each of these:

- Decision function

$$\hat{y} = f_{\theta}(\mathbf{x}_i)$$

- Loss function

$$\ell(\hat{y}, \mathbf{y}_i) \in \mathbb{R}$$

4. Train with SGD:  
(take small steps  
opposite the gradient)

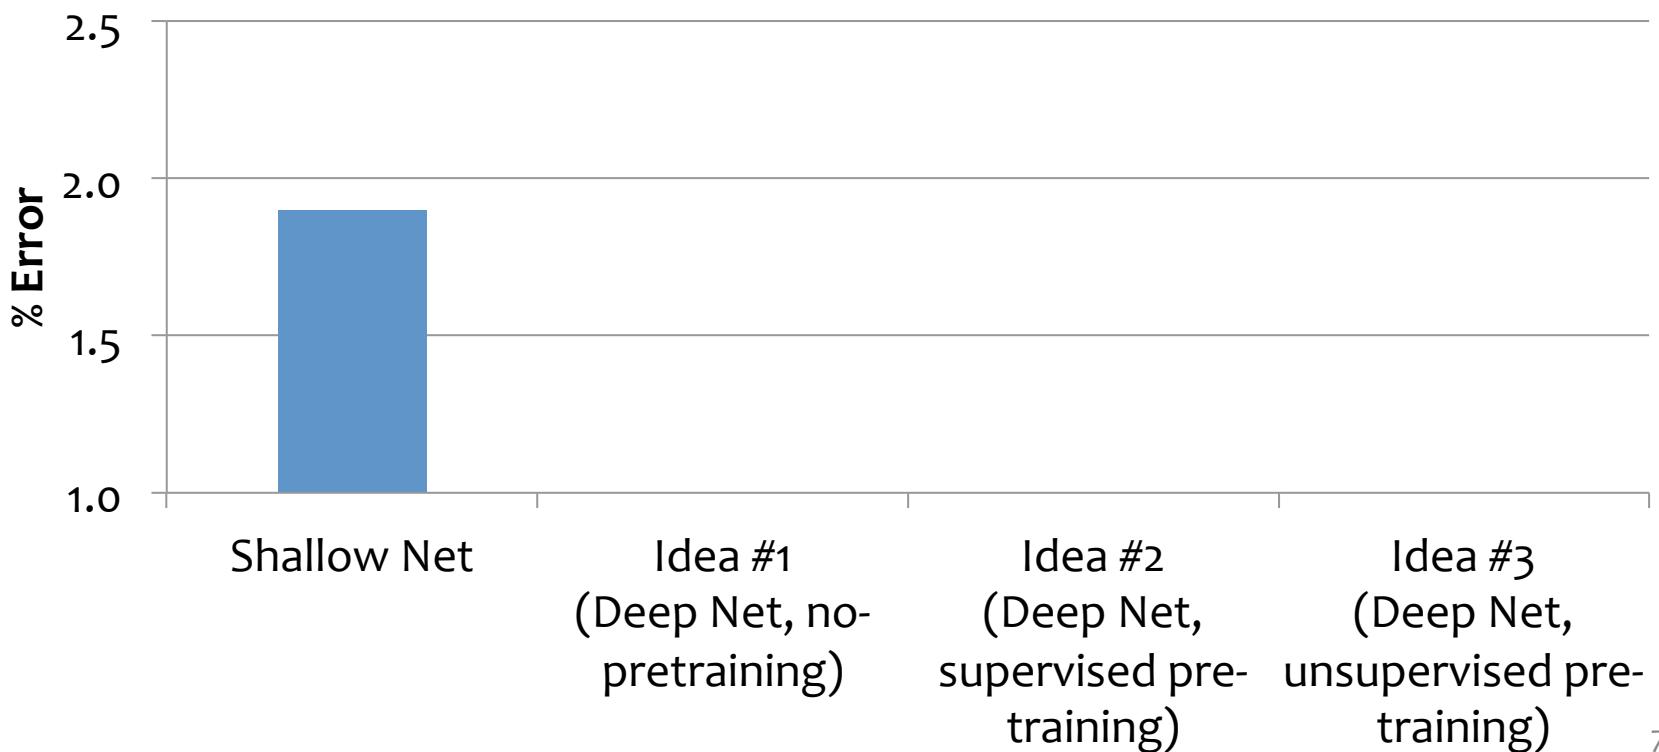
$$\theta^{(t+1)} = \theta^{(t)} - \eta_t \nabla \ell(f_{\theta}(\mathbf{x}_i), \mathbf{y}_i)$$

## Idea #1: No pre-training

- **Idea #1: (Just like a shallow network)**
  - Compute the supervised gradient by backpropagation.
  - Take small steps in the direction of the gradient (SGD)

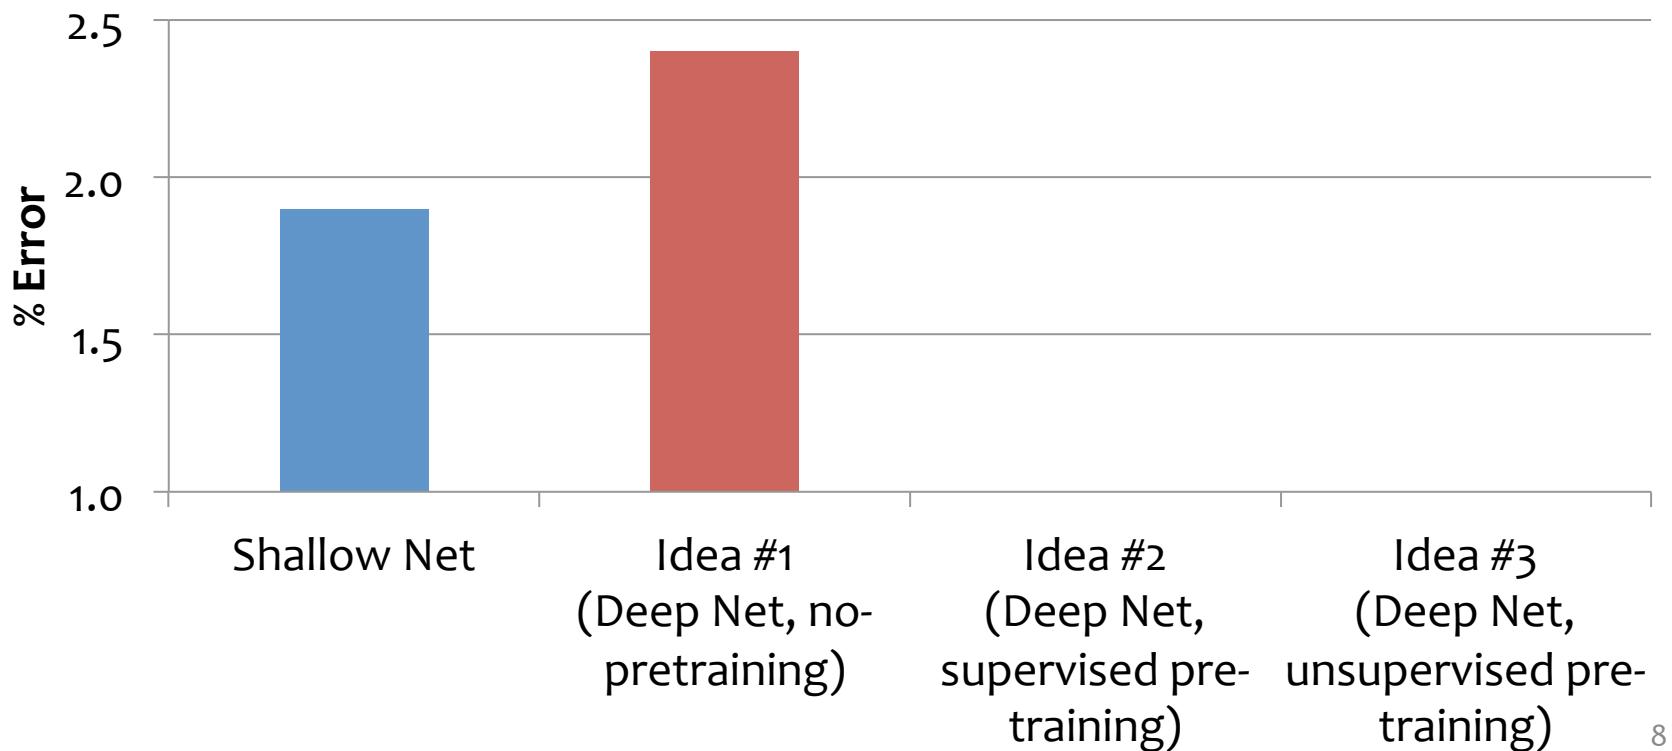
# Comparison on MNIST

- Results from Bengio et al. (2006) on MNIST digit classification task
- Percent error (lower is better)



# Comparison on MNIST

- Results from Bengio et al. (2006) on MNIST digit classification task
- Percent error (lower is better)

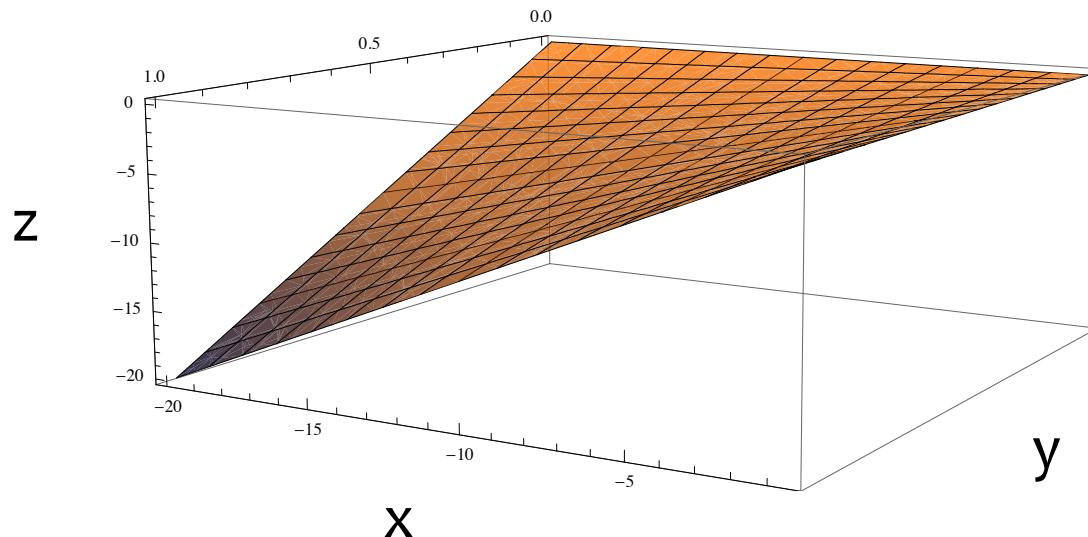


## Idea #1: No pre-training

- Idea #1: (Just like a shallow network)
  - Compute the supervised gradient by backpropagation.
  - Take small steps in the direction of the gradient (SGD)
- What goes wrong?
  - A. Gets stuck in local optima
    - Nonconvex objective
    - Usually start at a random (bad) point in parameter space
  - B. Gradient is progressively getting more dilute
    - “Vanishing gradients”

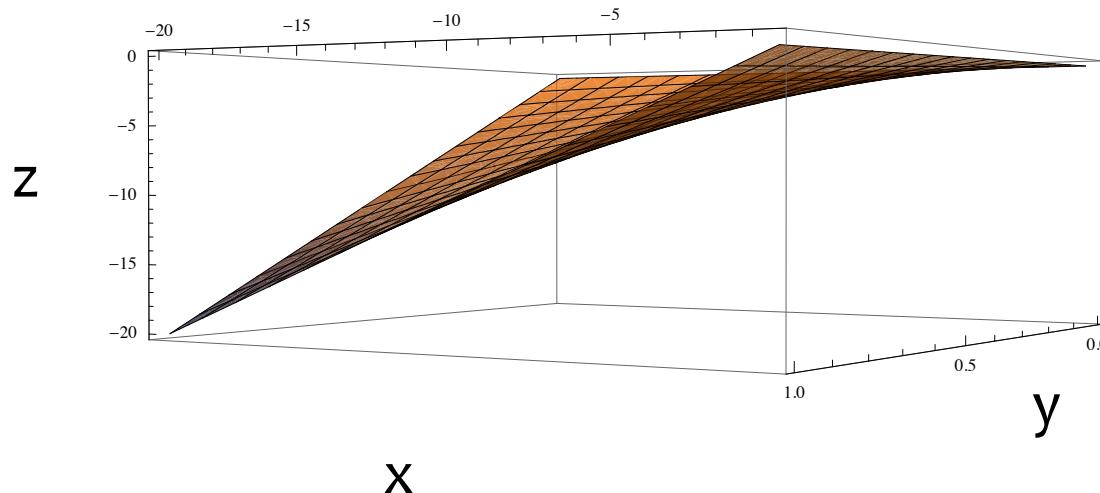
# Problem A: Nonconvexity

- Where does the nonconvexity come from?
- Even a simple quadratic  $z = xy$  objective is nonconvex:



# Problem A: Nonconvexity

- Where does the nonconvexity come from?
- Even a simple quadratic  $z = xy$  objective is nonconvex:

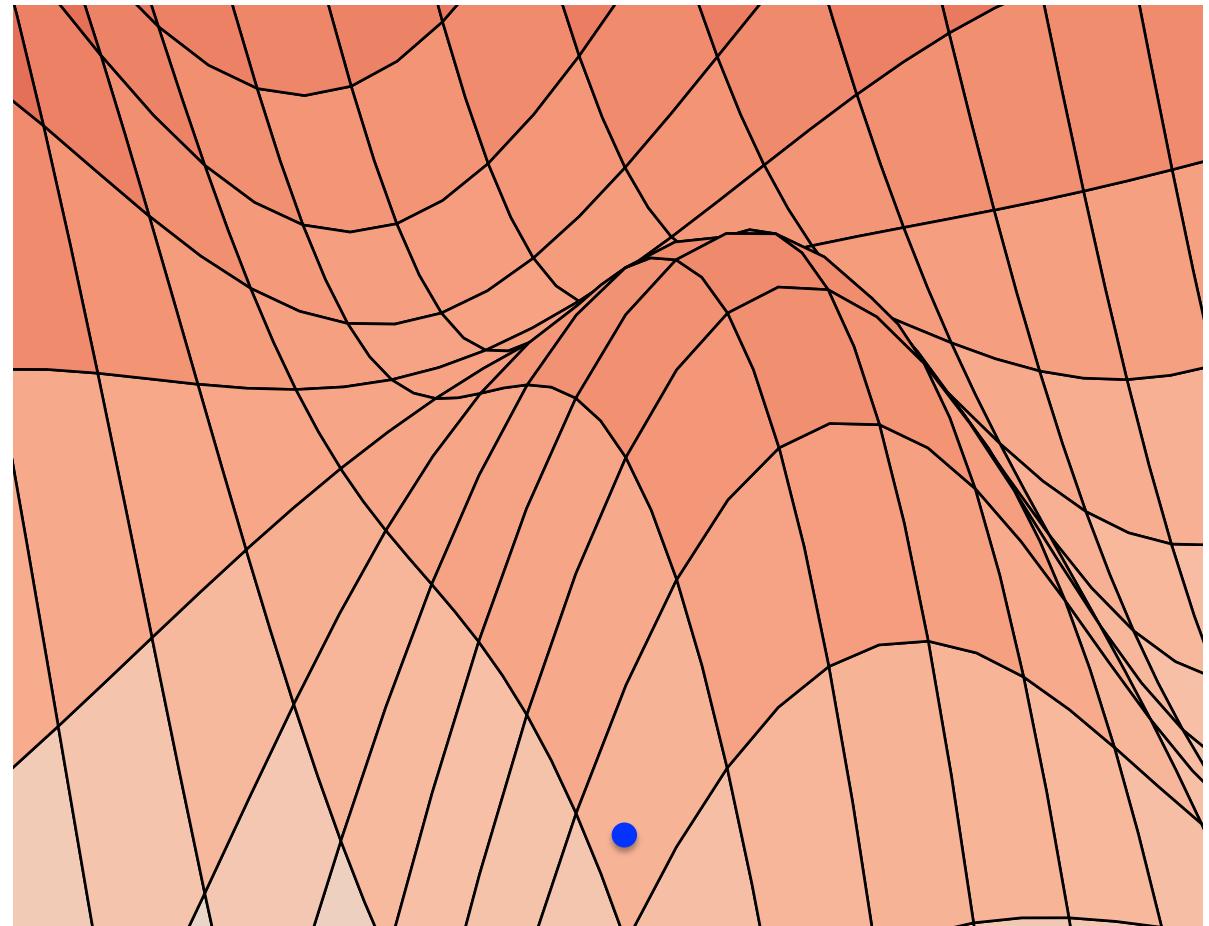


# Training

# Problem A: Nonconvexity

Stochastic Gradient  
Descent...

...climbs to the top  
of the nearest hill...

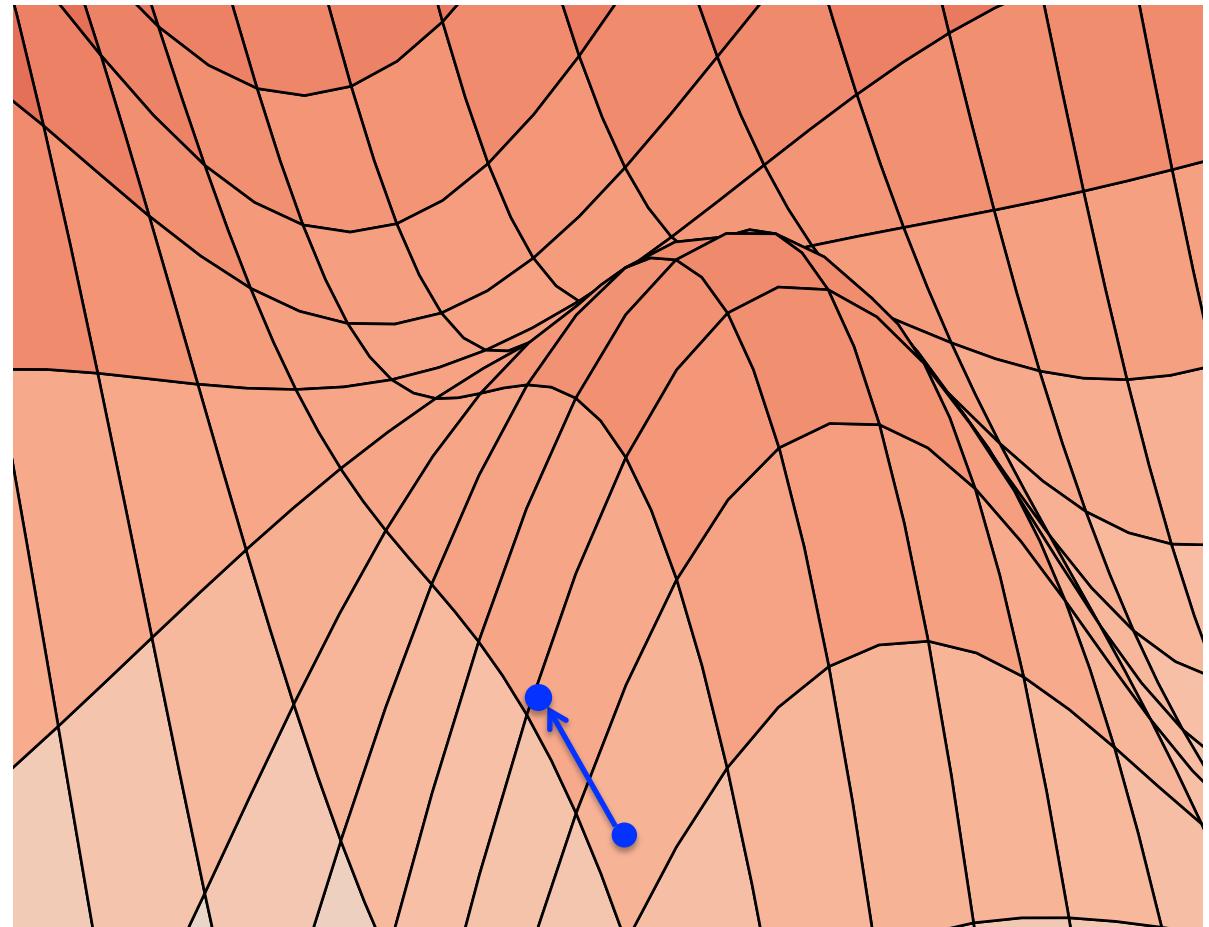


# Training

# Problem A: Nonconvexity

Stochastic Gradient  
Descent...

...climbs to the top  
of the nearest hill...

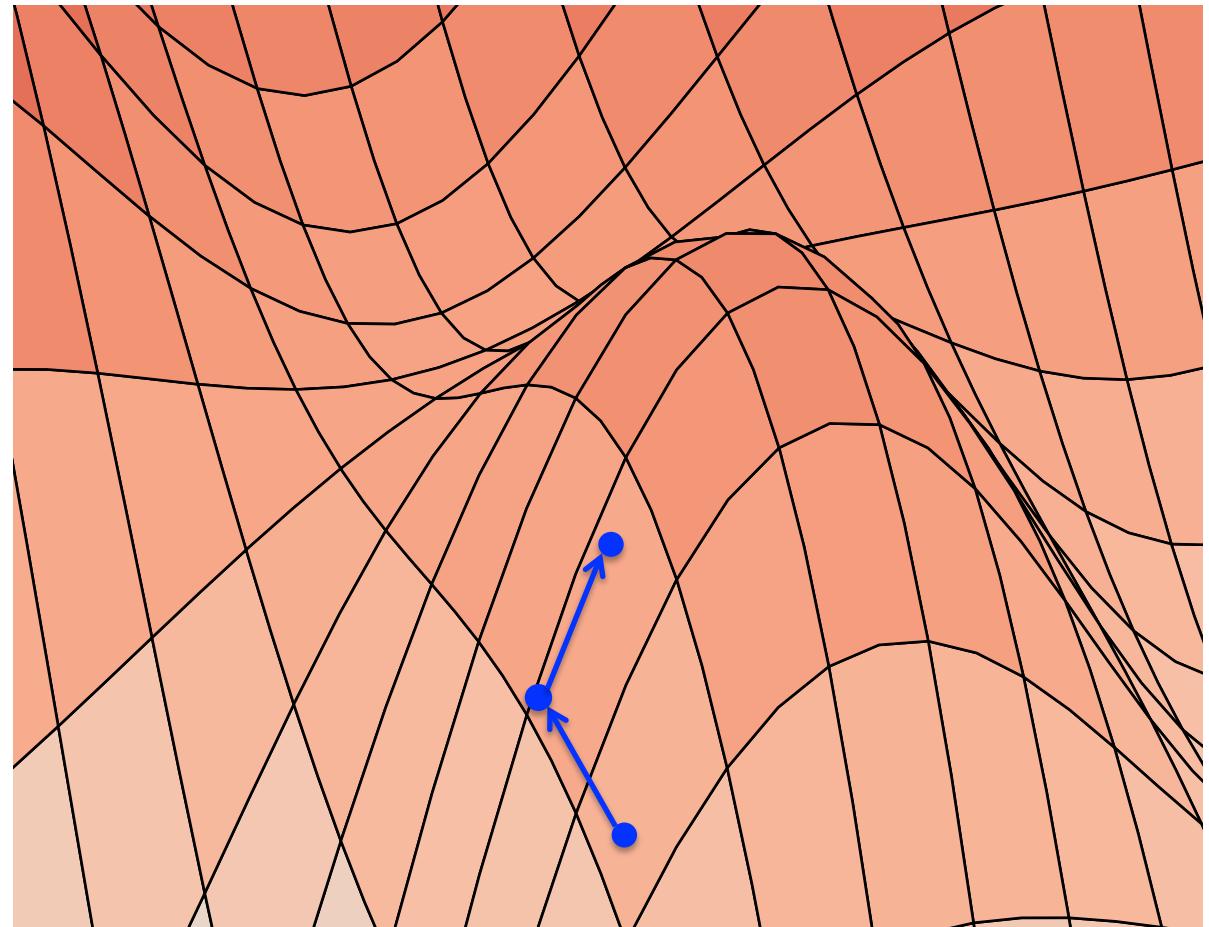


# Training

# Problem A: Nonconvexity

Stochastic Gradient  
Descent...

...climbs to the top  
of the nearest hill...

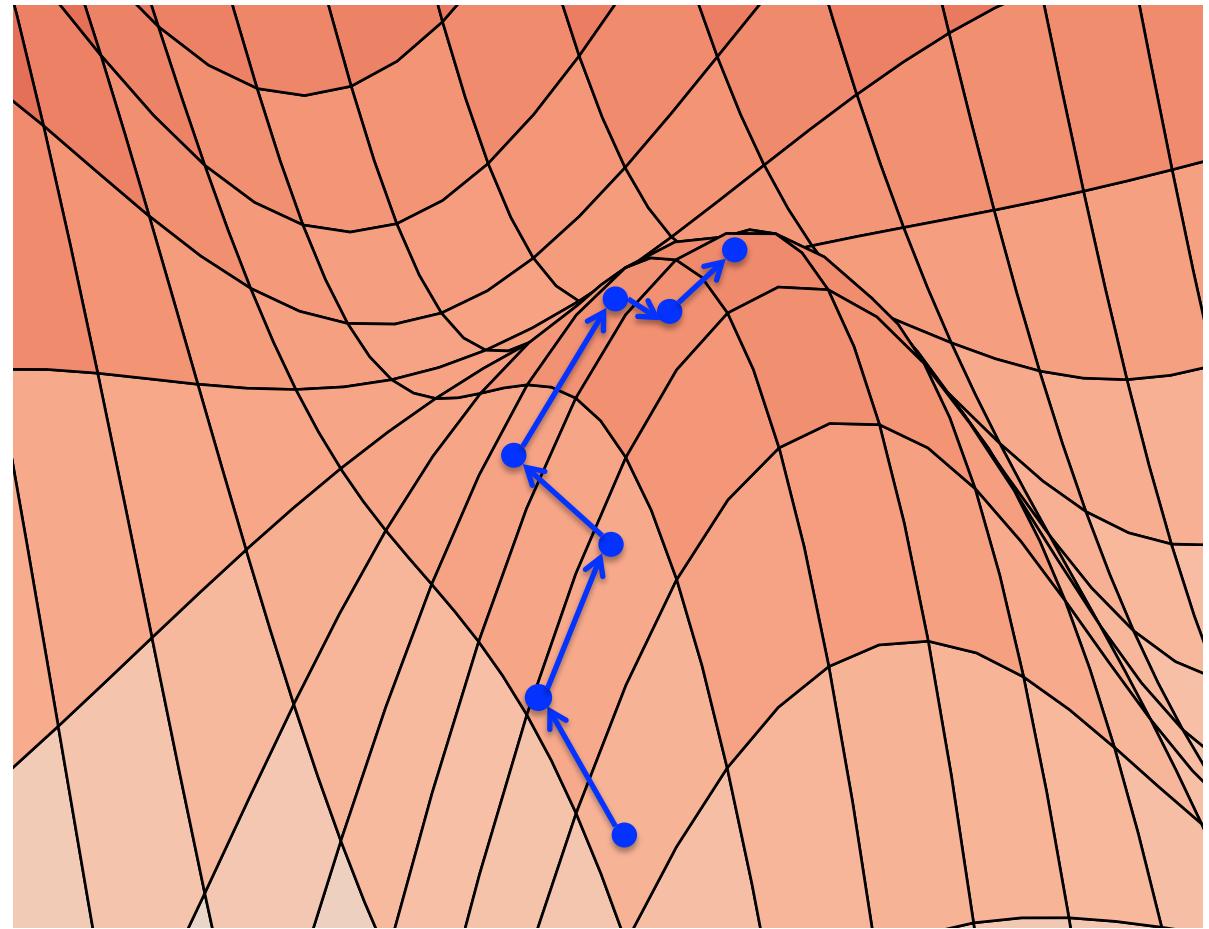


# Training

# Problem A: Nonconvexity

Stochastic Gradient  
Descent...

...climbs to the top  
of the nearest hill...



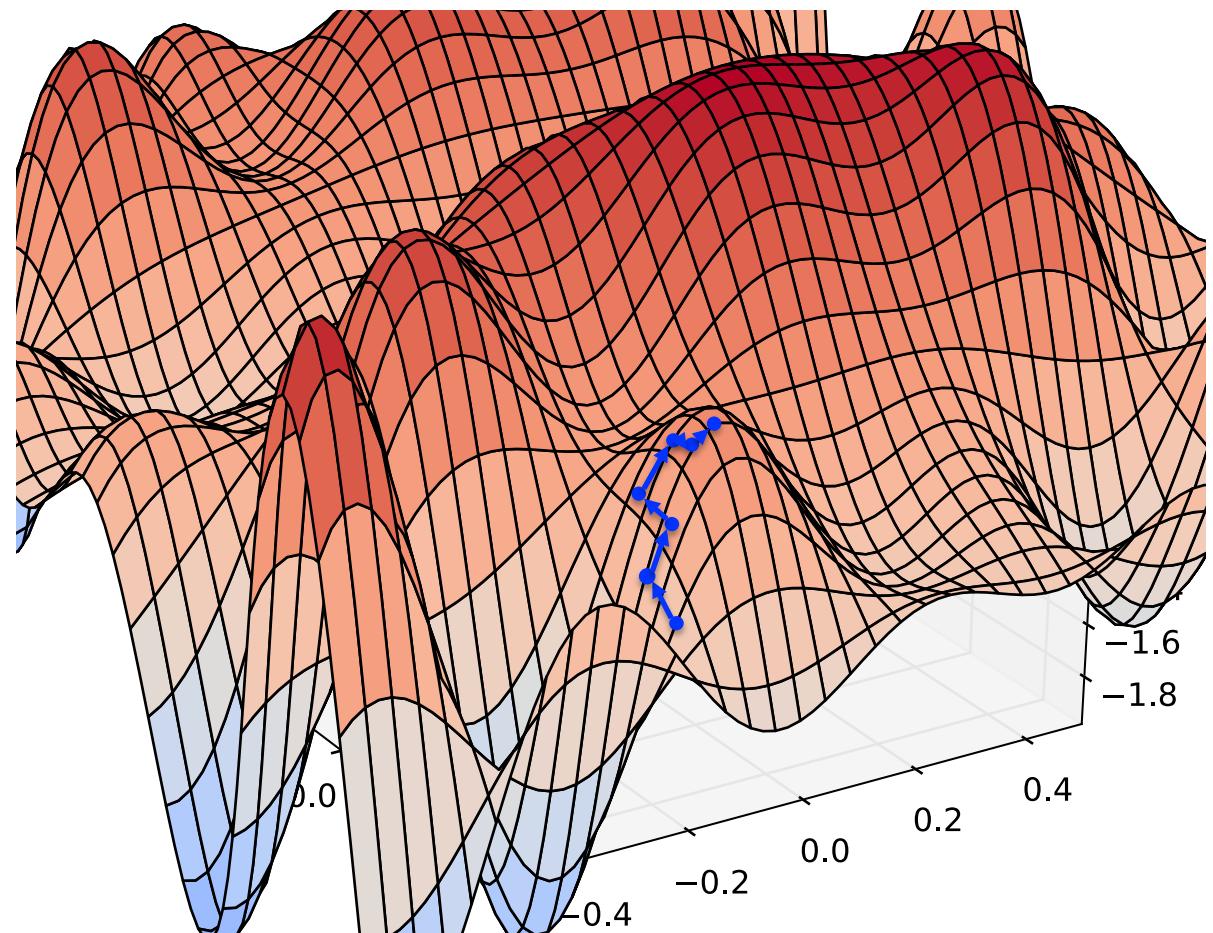
# Training

# Problem A: Nonconvexity

Stochastic Gradient  
Descent...

...climbs to the top  
of the nearest hill...

...which might not  
lead to the top of  
the mountain

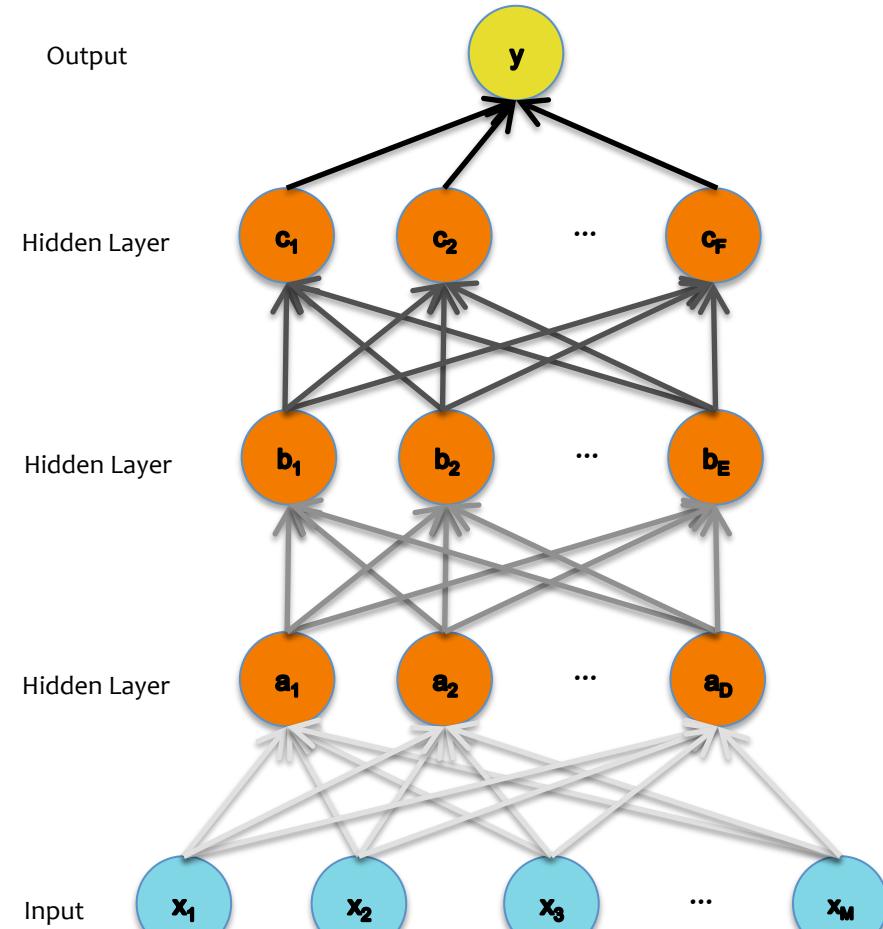


# Training

# Problem B: Vanishing Gradients

The gradient for an edge at the base of the network depends on the gradients of many edges above it

The chain rule multiplies many of these partial derivatives together

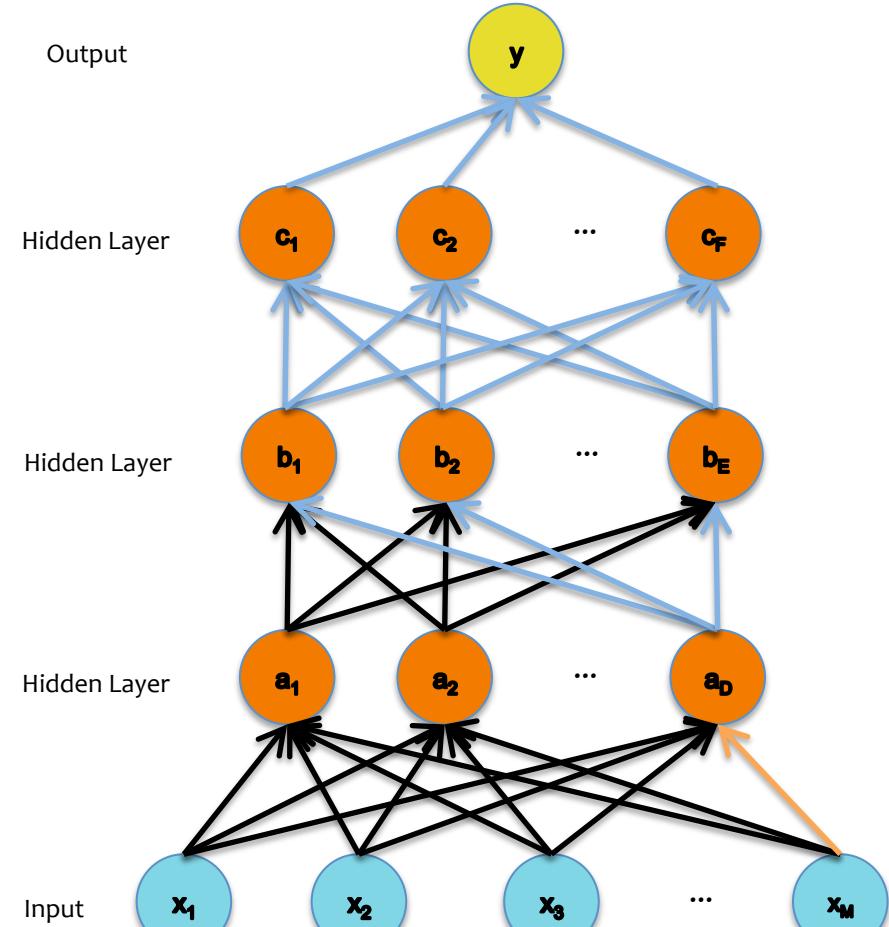


# Training

# Problem B: Vanishing Gradients

The gradient for an edge at the base of the network depends on the gradients of many edges above it

The chain rule multiplies many of these partial derivatives together

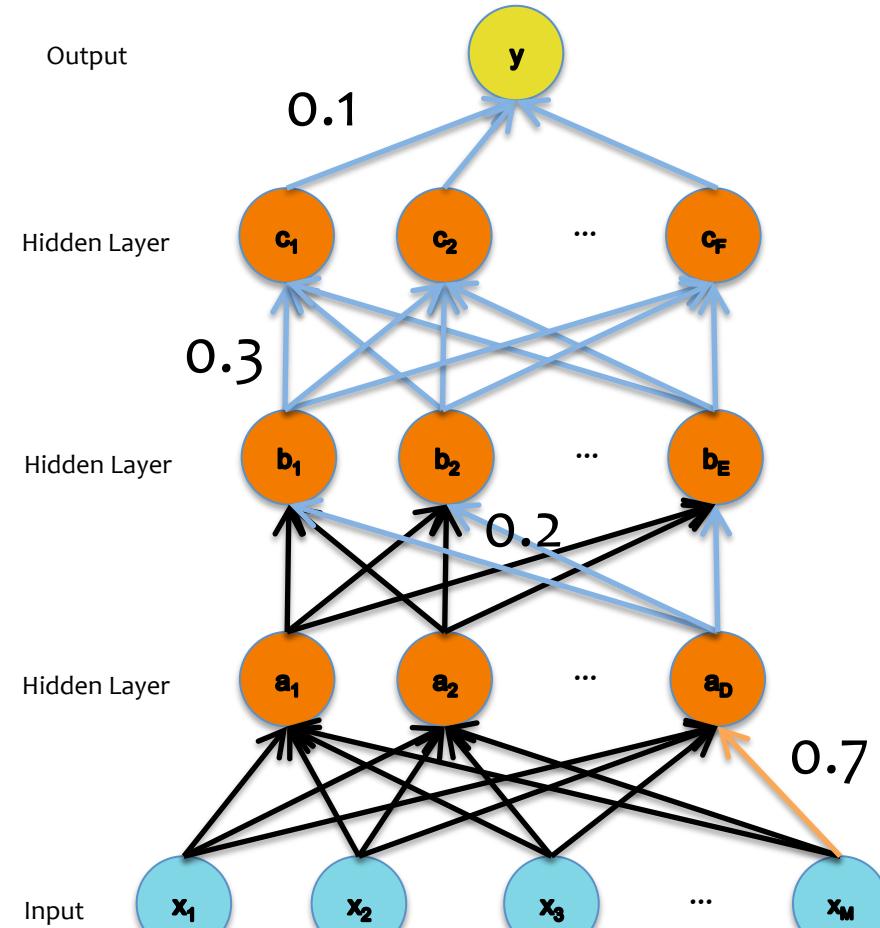


# Training

# Problem B: Vanishing Gradients

The gradient for an edge at the base of the network depends on the gradients of many edges above it

The chain rule multiplies many of these partial derivatives together



## Idea #1: No pre-training

- Idea #1: (Just like a shallow network)
  - Compute the supervised gradient by backpropagation.
  - Take small steps in the direction of the gradient (SGD)
- What goes wrong?
  - A. Gets stuck in local optima
    - Nonconvex objective
    - Usually start at a random (bad) point in parameter space
  - B. Gradient is progressively getting more dilute
    - “Vanishing gradients”

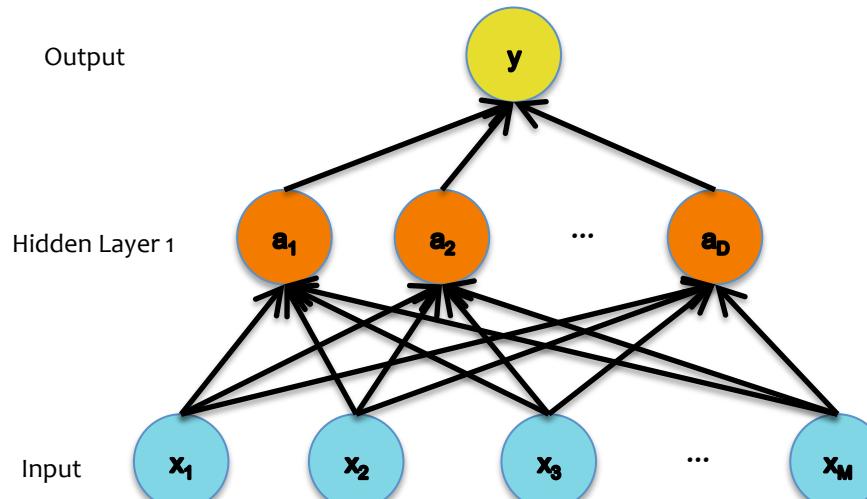
# Idea #2: Supervised Pre-training

- Idea #2: (Two Steps)
    - Train each level of the model in a **greedy** way
    - Then use our **original idea**
1. Supervised Pre-training
    - Use **labeled** data
    - Work bottom-up
      - Train hidden layer 1. Then fix its parameters.
      - Train hidden layer 2. Then fix its parameters.
      - ...
      - Train hidden layer n. Then fix its parameters.
  2. Supervised Fine-tuning
    - Use **labeled** data to train following “Idea #1”
    - Refine the features by backpropagation so that they become tuned to the end-task

# Training

# Idea #2: Supervised Pre-training

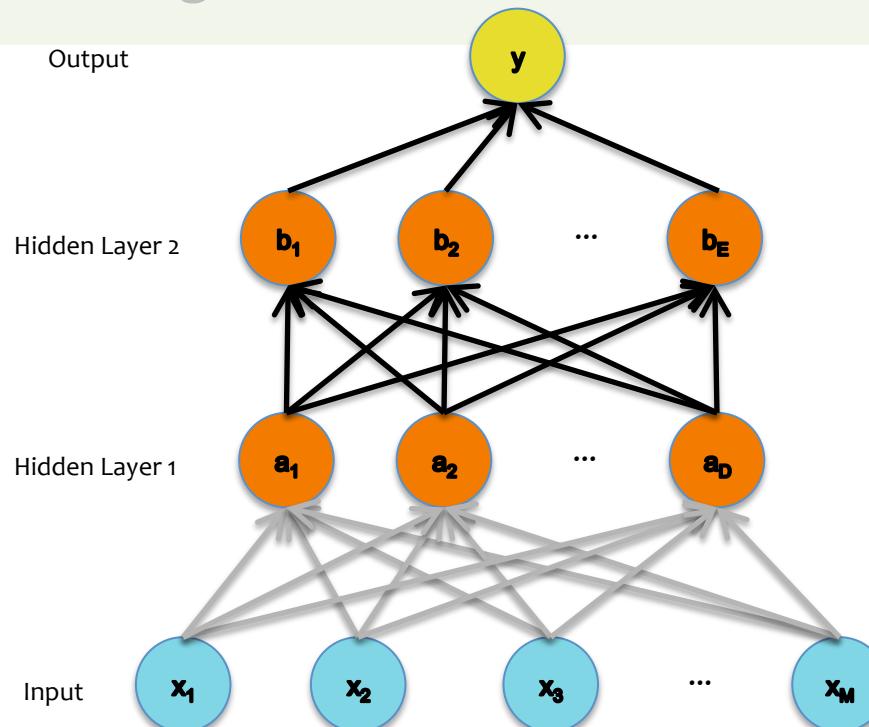
- Idea #2: (Two Steps)
  - Train each level of the model in a **greedy** way
  - Then use our **original idea**



# Training

# Idea #2: Supervised Pre-training

- Idea #2: (Two Steps)
  - Train each level of the model in a *greedy* way
  - Then use our **original idea**

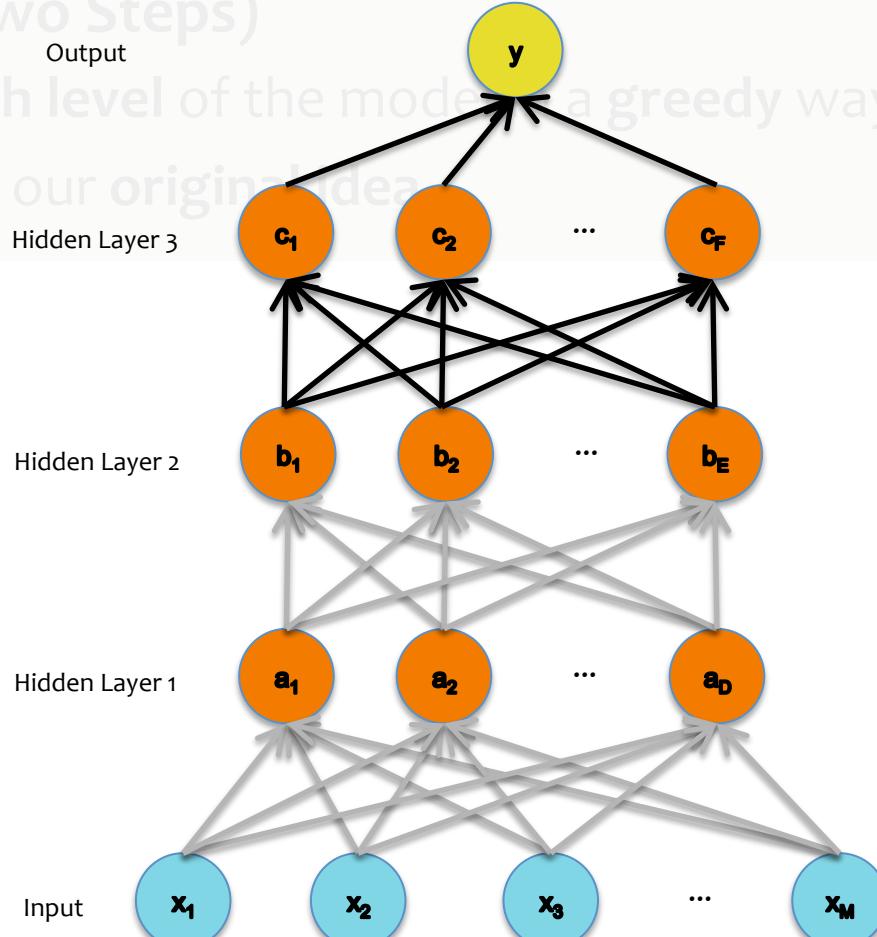


# Training

# Idea #2: Supervised Pre-training

- Idea #2: (Two Steps)

- Train each level of the model in a greedy way
- Then use our original idea

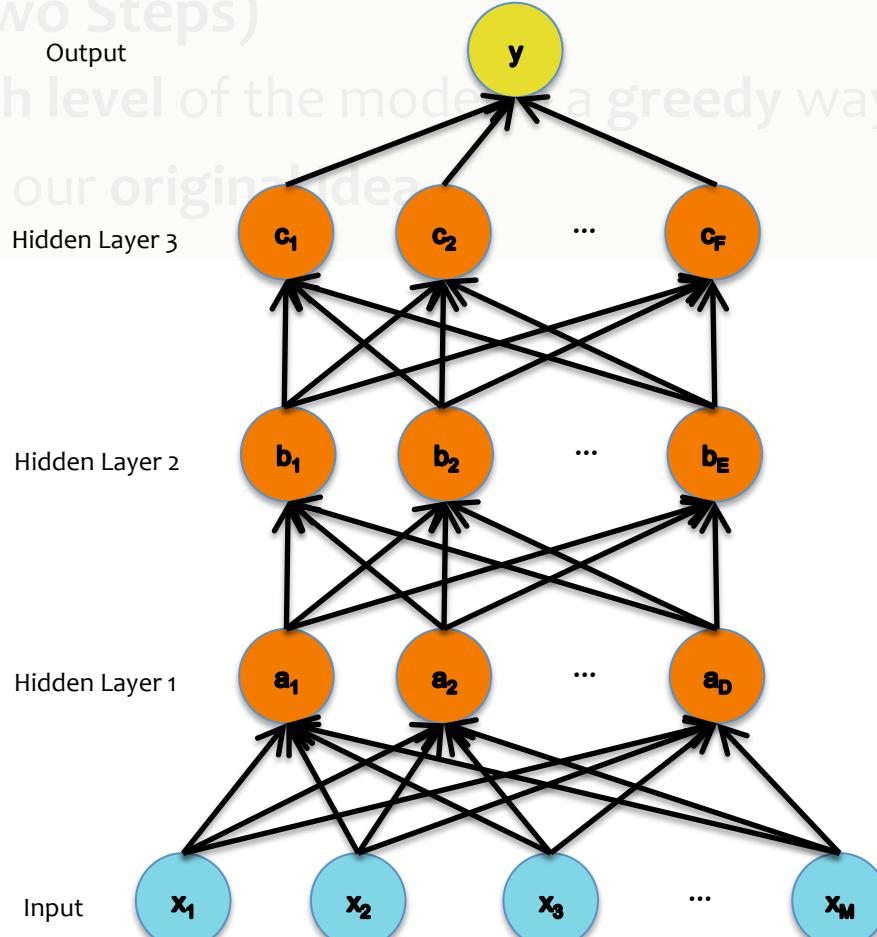


# Training

# Idea #2: Supervised Pre-training

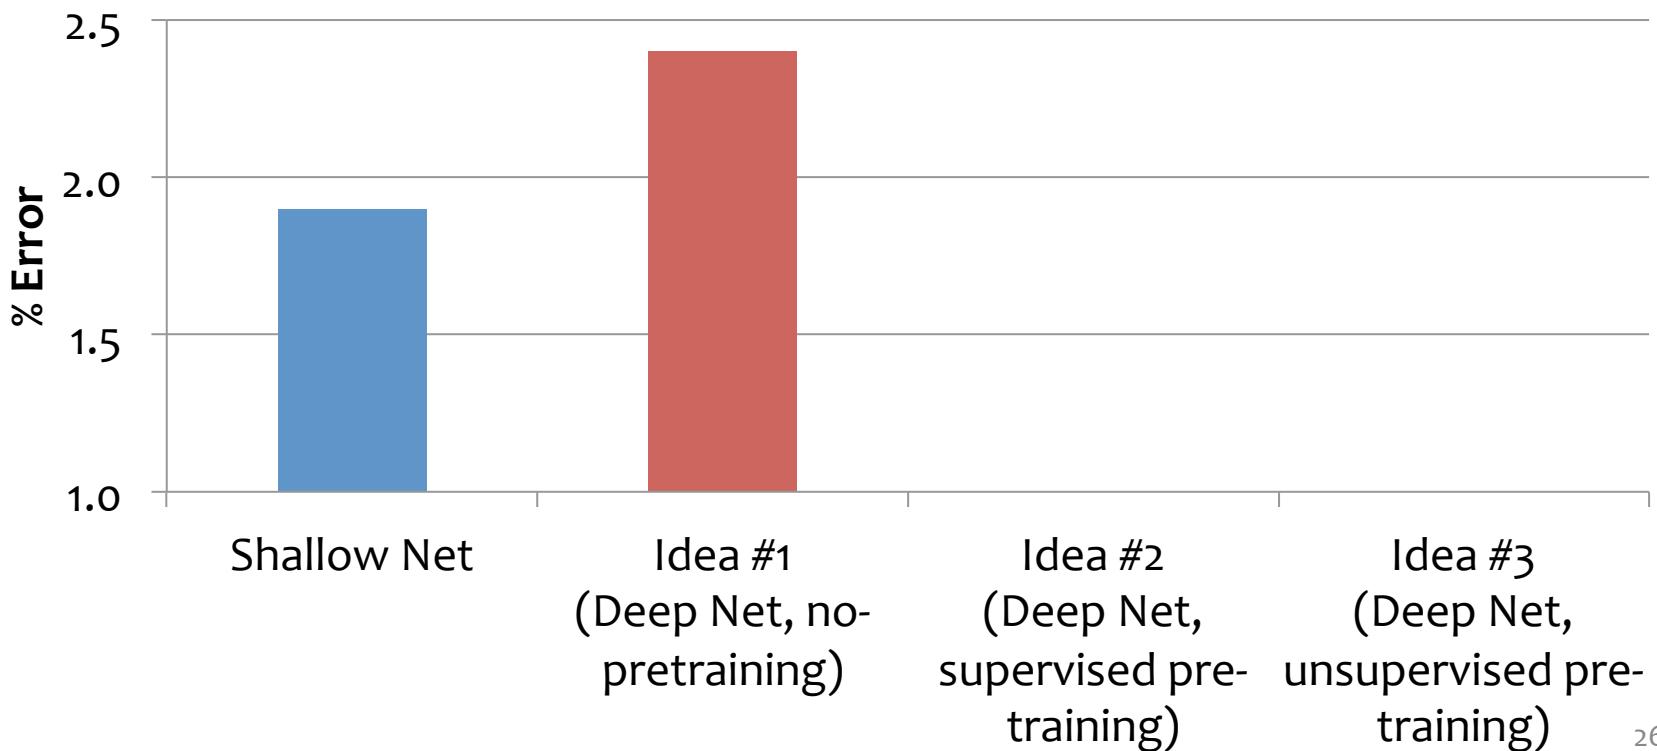
- Idea #2: (Two Steps)

- Train each level of the model in a greedy way
- Then use our original idea



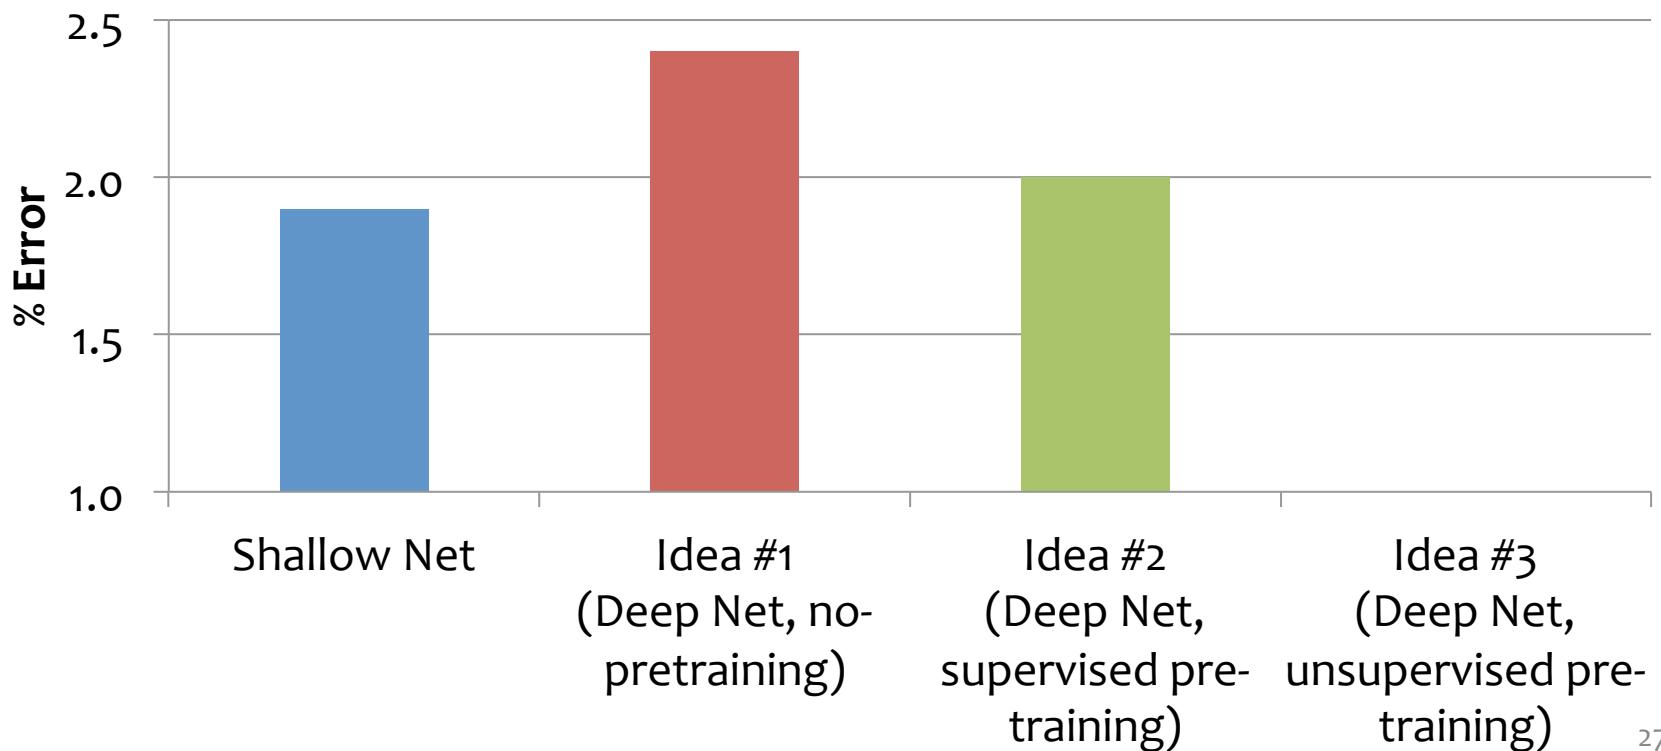
# Comparison on MNIST

- Results from Bengio et al. (2006) on MNIST digit classification task
- Percent error (lower is better)



# Comparison on MNIST

- Results from Bengio et al. (2006) on MNIST digit classification task
- Percent error (lower is better)



# Idea #3: Unsupervised Pre-training

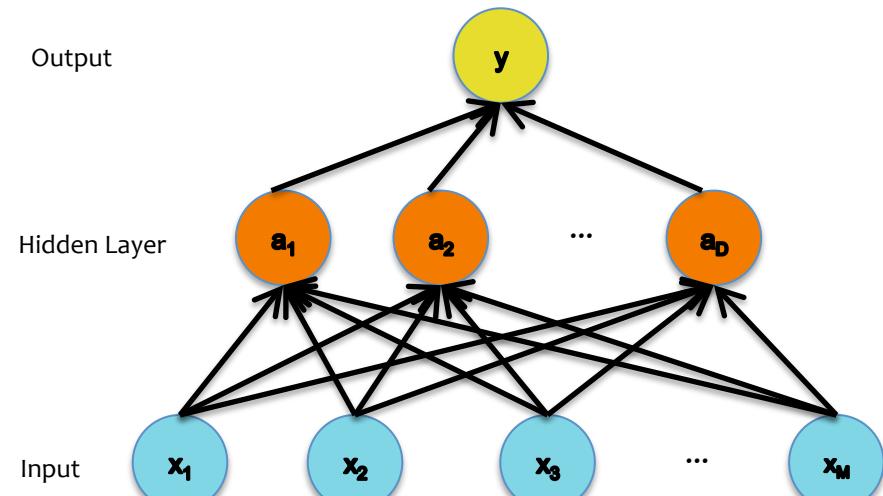
- Idea #3: (Two Steps)
  - Use our original idea, but **pick a better starting point**
  - **Train each level** of the model in a **greedy** way

1. Unsupervised Pre-training
  - Use **unlabeled** data
  - Work bottom-up
    - Train hidden layer 1. Then fix its parameters.
    - Train hidden layer 2. Then fix its parameters.
    - ...
    - Train hidden layer n. Then fix its parameters.
2. Supervised Fine-tuning
  - Use **labeled** data to train following “Idea #1”
  - Refine the features by backpropagation so that they become tuned to the end-task

# The solution: Unsupervised pre-training

## Unsupervised pre-training of the first layer:

- What should it predict?
- What else do we observe?
- **The input!**

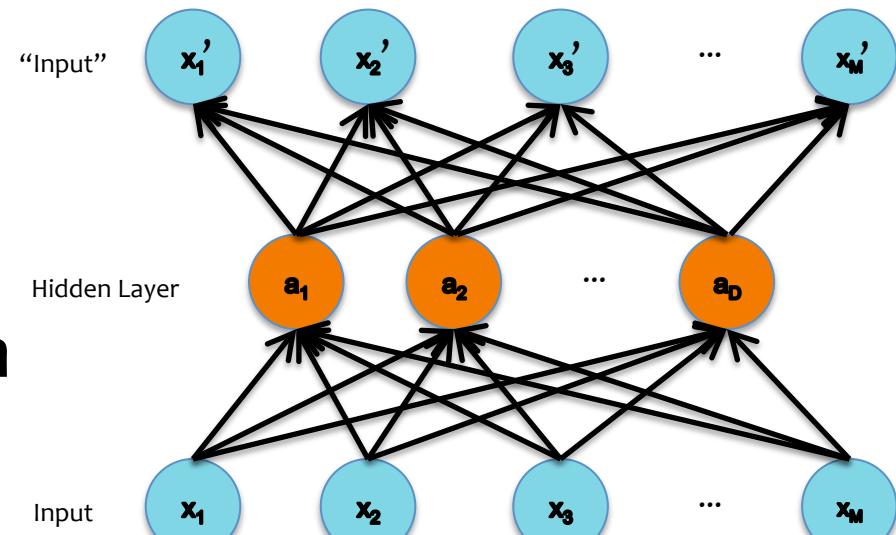


# The solution: Unsupervised pre-training

## Unsupervised pre-training of the first layer:

- What should it predict?
- What else do we observe?
- **The input!**

This topology defines an Auto-encoder.



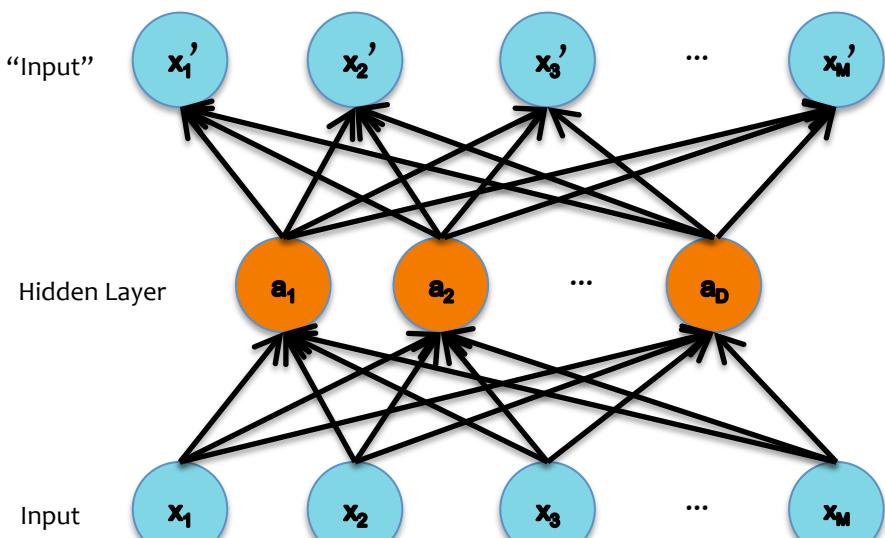
# Auto-Encoders

Key idea: Encourage  $z$  to give small reconstruction error:

- $x'$  is the reconstruction of  $x$
- Loss =  $\| x - \text{DECODER}(\text{ENCODER}(x)) \|_2^2$
- Train with the same backpropagation algorithm for 2-layer Neural Networks with  $x_m$  as both input and output.

DECODER:  $x' = h(W'z)$

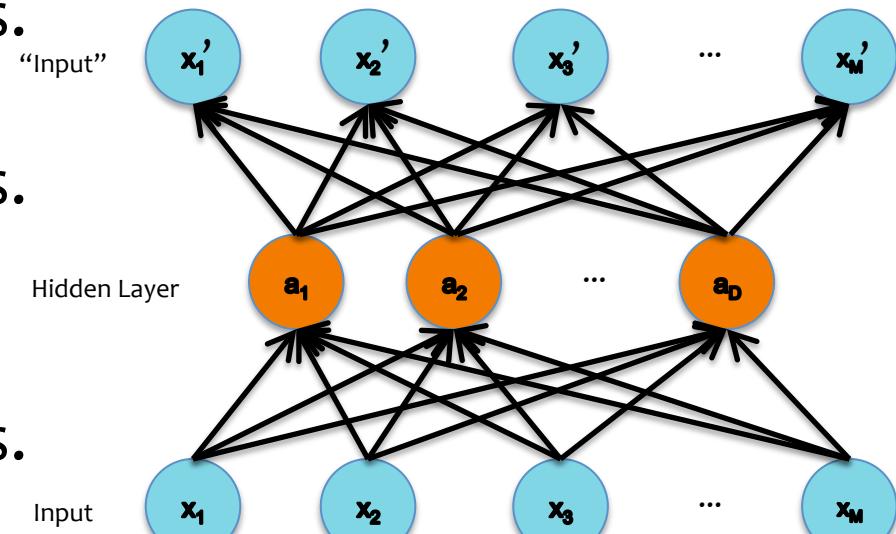
ENCODER:  $z = h(Wx)$



# The solution: Unsupervised pre-training

## Unsupervised pre-training

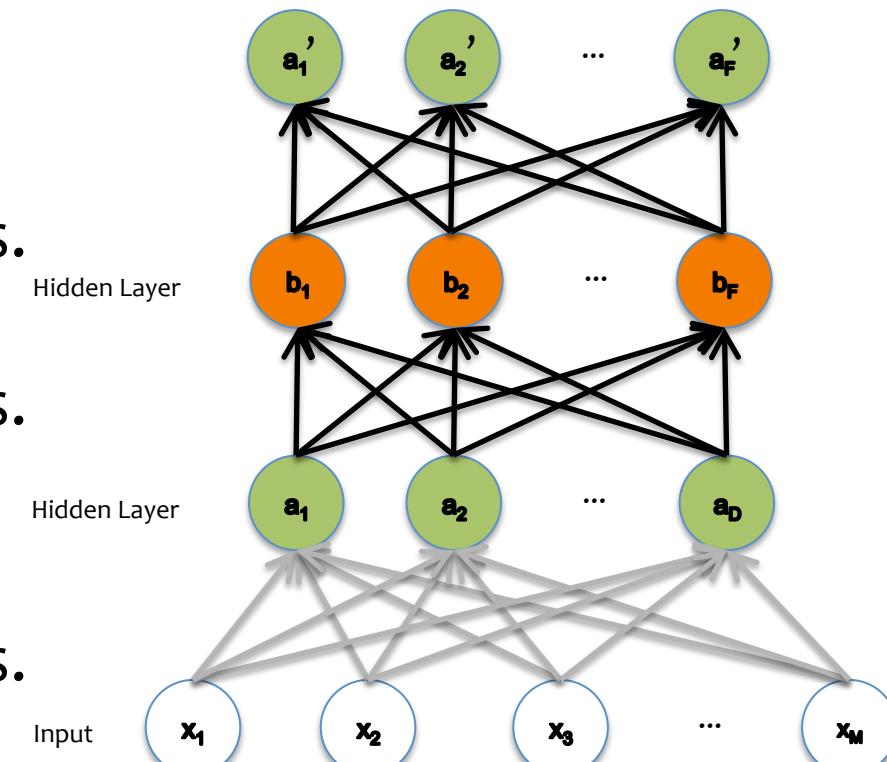
- Work bottom-up
  - Train hidden layer 1.  
Then fix its parameters.
  - Train hidden layer 2.  
Then fix its parameters.
  - ...
  - Train hidden layer n.  
Then fix its parameters.



# The solution: Unsupervised pre-training

## Unsupervised pre-training

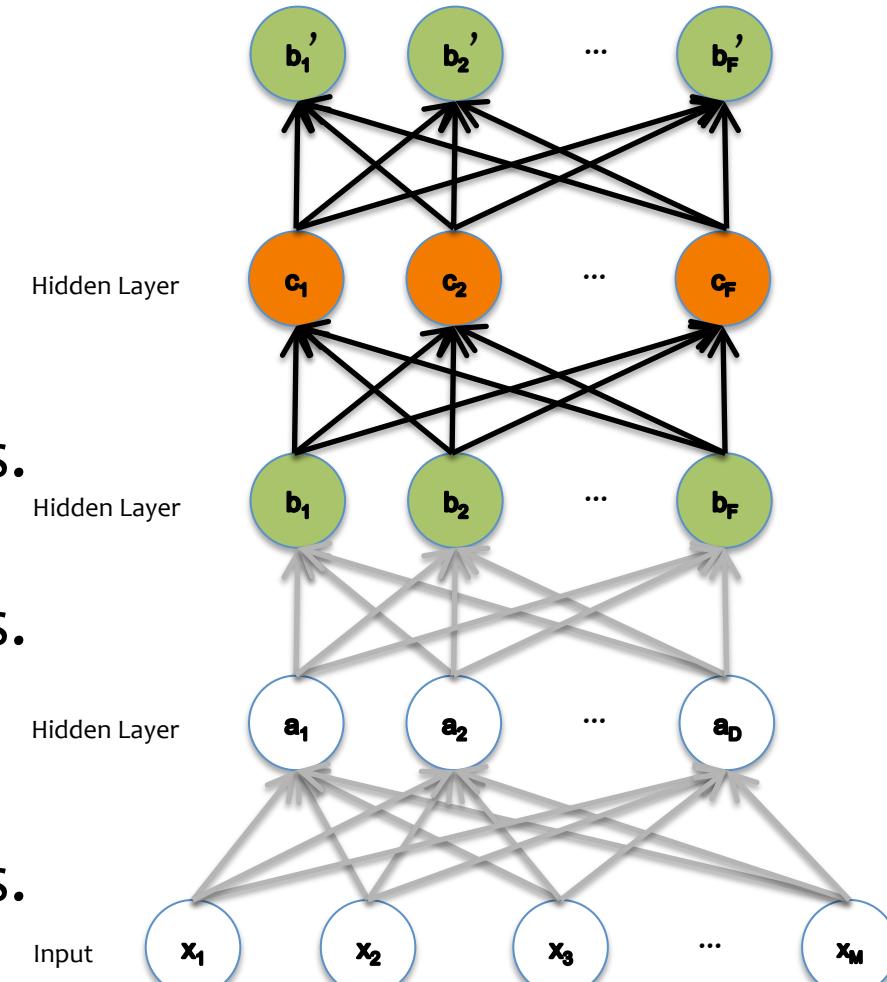
- Work bottom-up
  - Train hidden layer 1.  
Then fix its parameters.
  - Train hidden layer 2.  
Then fix its parameters.
  - ...
  - Train hidden layer n.  
Then fix its parameters.



# The solution: Unsupervised pre-training

## Unsupervised pre-training

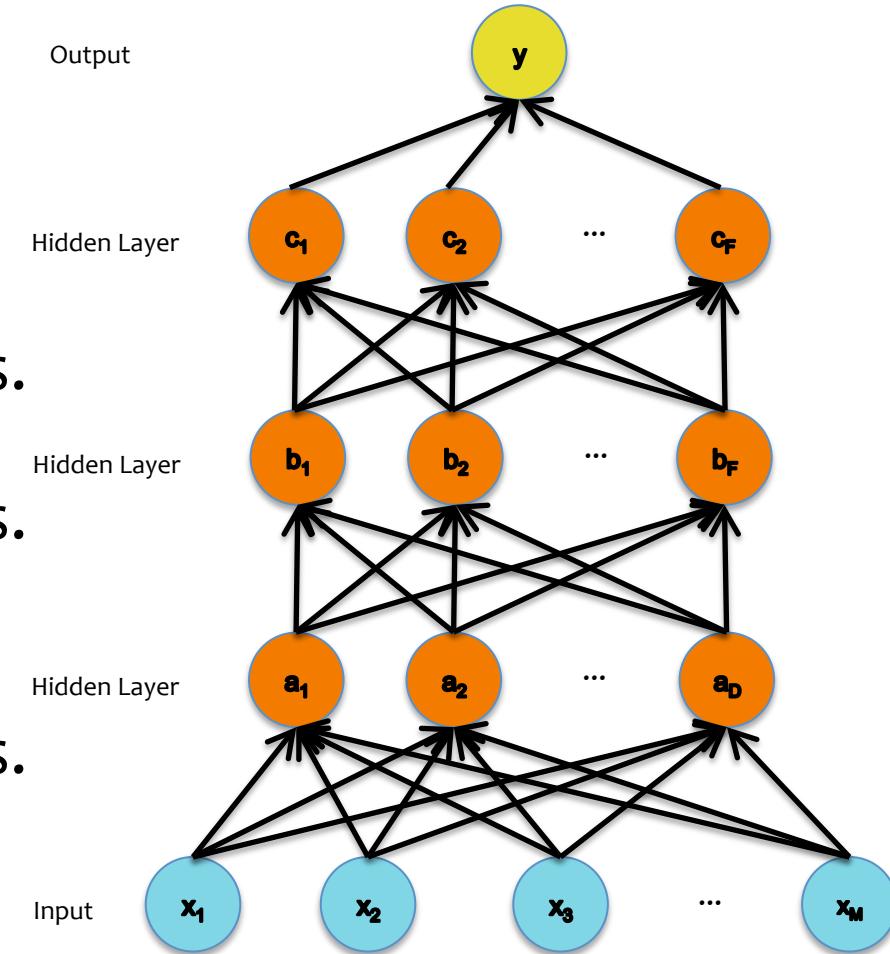
- Work bottom-up
  - Train hidden layer 1.  
Then fix its parameters.
  - Train hidden layer 2.  
Then fix its parameters.
  - ...
  - Train hidden layer n.  
Then fix its parameters.



# The solution: Unsupervised pre-training

## Unsupervised pre-training

- Work bottom-up
  - Train hidden layer 1. Then fix its parameters.
  - Train hidden layer 2. Then fix its parameters.
  - ...
  - Train hidden layer n. Then fix its parameters.



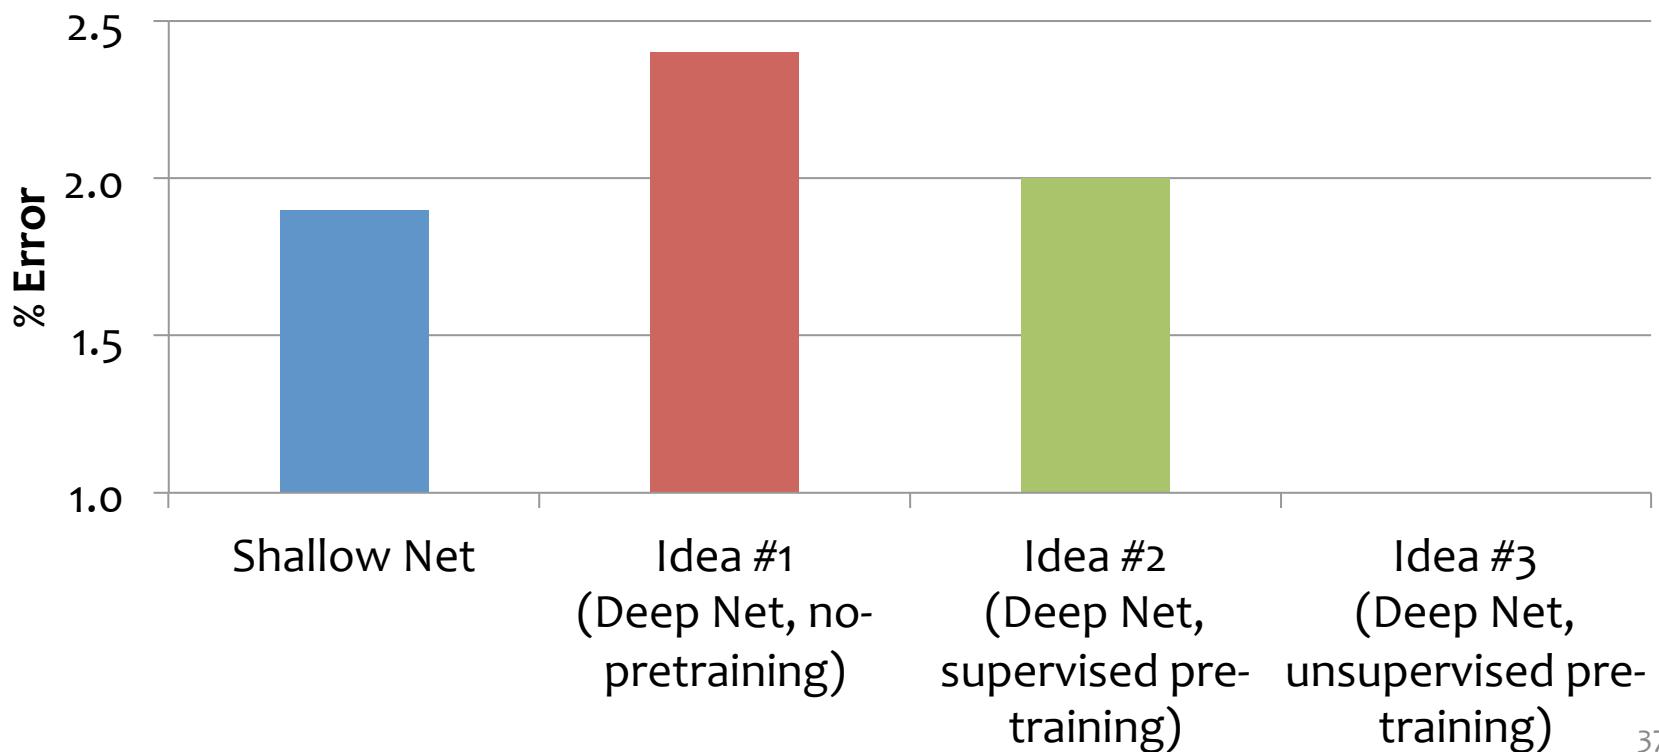
**Supervised fine-tuning**  
Backprop and update all  
parameters

# Deep Network Training

- **Idea #1:**
  1. Supervised fine-tuning only
- **Idea #2:**
  1. Supervised layer-wise pre-training
  2. Supervised fine-tuning
- **Idea #3:**
  1. Unsupervised layer-wise pre-training
  2. Supervised fine-tuning

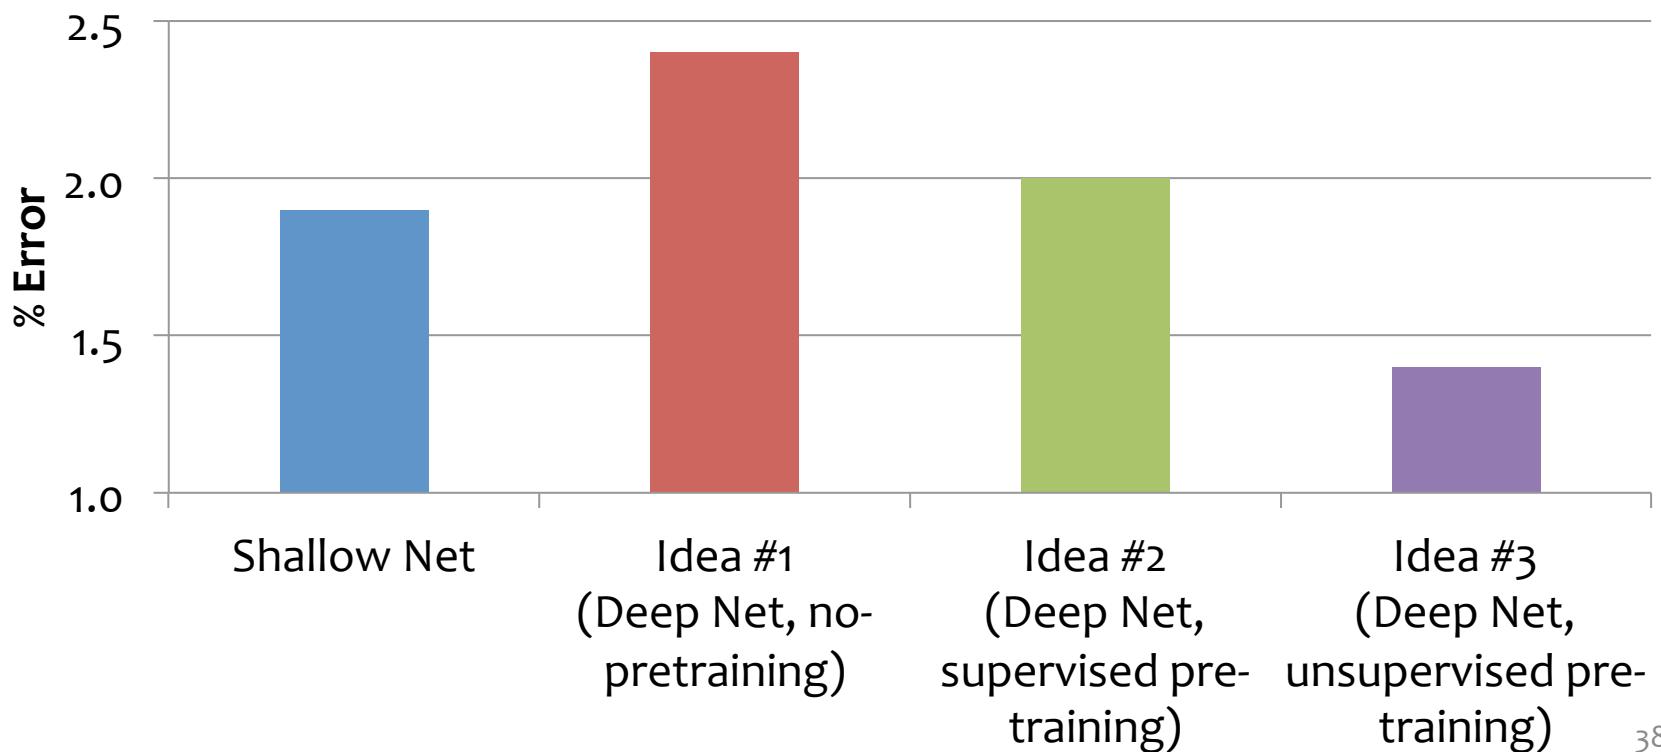
# Comparison on MNIST

- Results from Bengio et al. (2006) on MNIST digit classification task
- Percent error (lower is better)



# Comparison on MNIST

- Results from Bengio et al. (2006) on MNIST digit classification task
- Percent error (lower is better)



Training

# Is layer-wise pre-training always necessary?

**In 2010, a record on a hand-writing  
recognition task was set by standard supervised  
backpropagation (our Idea #1).**

**How?** A very fast implementation on GPUs.

See Ciresen et al. (2010)

# Deep Learning

- Goal: learn features at different levels of abstraction
- Training can be tricky due to...
  - Nonconvexity
  - Vanishing gradients
- Unsupervised layer-wise pre-training can help with both!

# **RECURRENT NEURAL NETWORKS**

# Recurrent Neural Networks (RNNs)

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units:  $\mathbf{h} = (h_1, h_2, \dots, h_T), h_i \in \mathcal{R}^J$

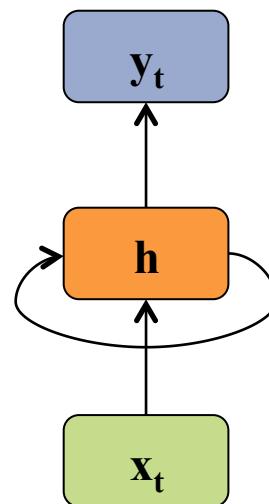
outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

nonlinearity:  $\mathcal{H}$

Definition of the RNN:

$$h_t = \mathcal{H}(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

$$y_t = W_{hy}h_t + b_y$$



# Recurrent Neural Networks (RNNs)

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units:  $\mathbf{h} = (h_1, h_2, \dots, h_T), h_i \in \mathcal{R}^J$

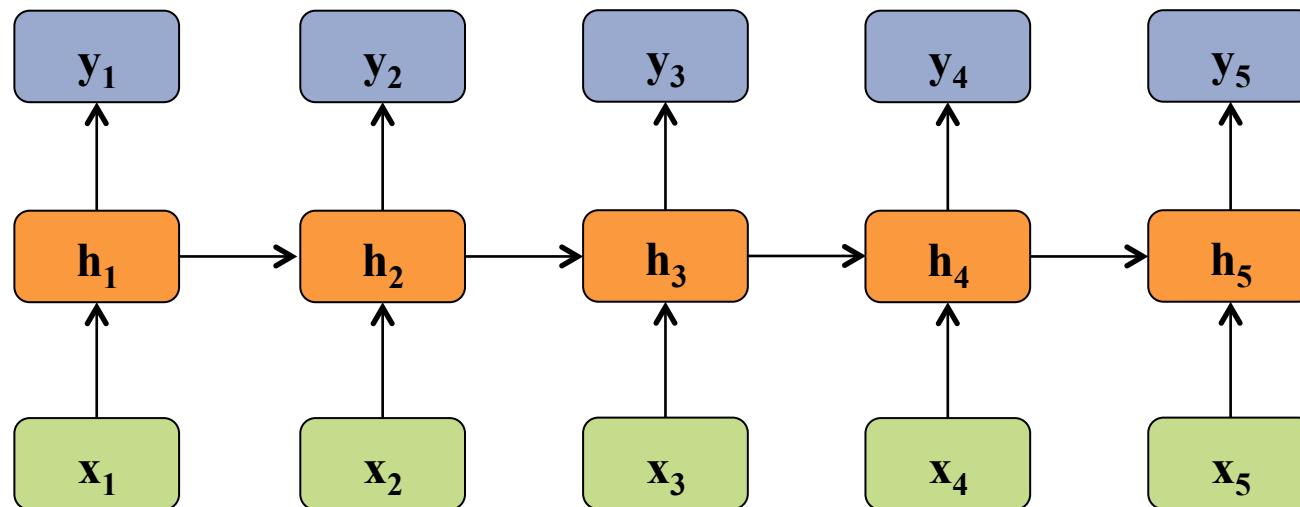
outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

nonlinearity:  $\mathcal{H}$

Definition of the RNN:

$$h_t = \mathcal{H}(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

$$y_t = W_{hy}h_t + b_y$$



# Recurrent Neural Networks (RNNs)

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units:  $\mathbf{h} = (h_1, h_2, \dots, h_T), h_i \in \mathcal{R}^J$

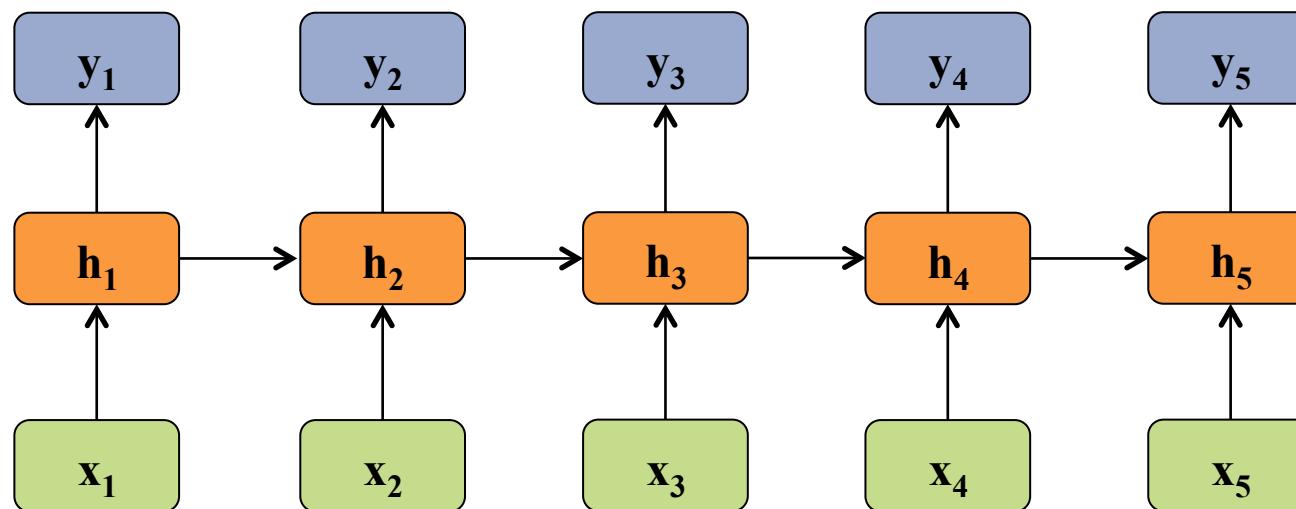
outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

nonlinearity:  $\mathcal{H}$

Definition of the RNN:

$$h_t = \mathcal{H}(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

$$y_t = W_{hy}h_t + b_y$$



# Recurrent Neural Networks (RNNs)

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units:  $\mathbf{h} = (h_1, h_2, \dots, h_T), h_i \in \mathcal{R}^J$

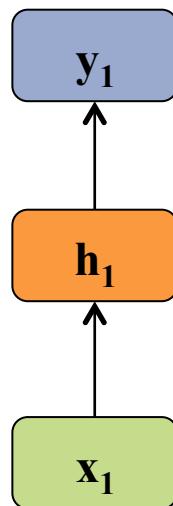
outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

nonlinearity:  $\mathcal{H}$

Definition of the RNN:

$$h_t = \mathcal{H}(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

$$y_t = W_{hy}h_t + b_y$$



- If  $T=1$ , then we have a standard feed-forward **neural net with one hidden layer**
- All of the deep nets from last lecture (DNN, DBN, DBM) required **fixed size inputs/outputs**

# Recurrent Neural Networks (RNNs)

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units:  $\mathbf{h} = (h_1, h_2, \dots, h_T), h_i \in \mathcal{R}^J$

outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

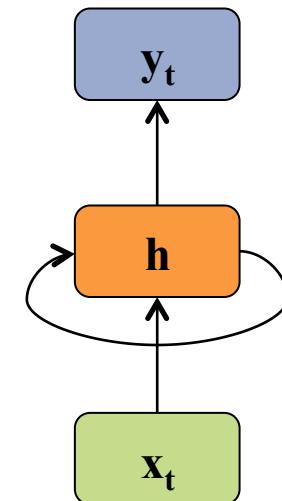
nonlinearity:  $\mathcal{H}$

Definition of the RNN:

$$h_t = \mathcal{H}(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

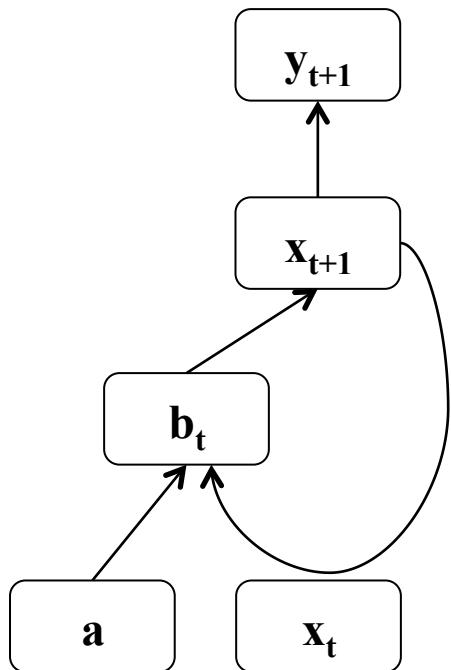
$$y_t = W_{hy}h_t + b_y$$

- By unrolling the RNN through time, we can **share parameters** and accommodate **arbitrary length** input/output pairs
- Applications: **time-series data** such as sentences, speech, stock-market, signal data, etc.



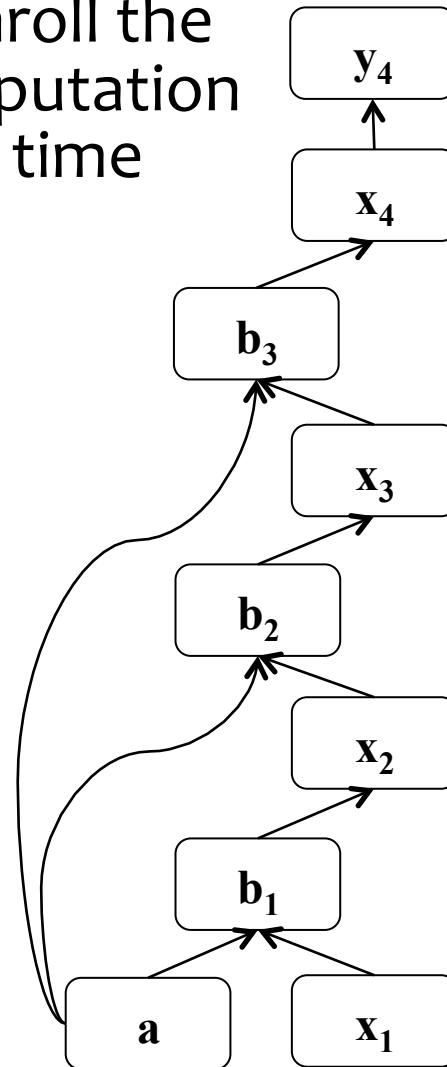
# Background: Backprop through time

**Recurrent neural network:**



**BPTT:**

1. Unroll the computation over time



2. Run backprop through the resulting feed-forward network

(Robinson & Fallside, 1981)  
(Werbos, 1988)  
(Mozer, 1995)

# Bidirectional RNN

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units:  $\overrightarrow{\mathbf{h}}$  and  $\overleftarrow{\mathbf{h}}$

outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

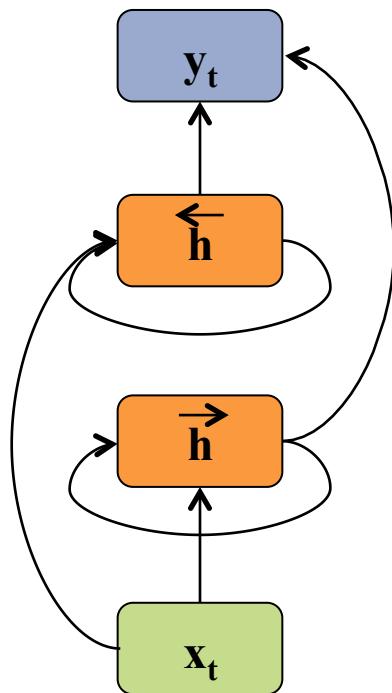
nonlinearity:  $\mathcal{H}$

Recursive Definition:

$$\overrightarrow{\mathbf{h}}_t = \mathcal{H} \left( W_{x \overrightarrow{\mathbf{h}}} x_t + W_{\overrightarrow{\mathbf{h}} \overrightarrow{\mathbf{h}}} \overrightarrow{\mathbf{h}}_{t-1} + b_{\overrightarrow{\mathbf{h}}} \right)$$

$$\overleftarrow{\mathbf{h}}_t = \mathcal{H} \left( W_{x \overleftarrow{\mathbf{h}}} x_t + W_{\overleftarrow{\mathbf{h}} \overleftarrow{\mathbf{h}}} \overleftarrow{\mathbf{h}}_{t+1} + b_{\overleftarrow{\mathbf{h}}} \right)$$

$$y_t = W_{\overrightarrow{\mathbf{h}} y} \overrightarrow{\mathbf{h}}_t + W_{\overleftarrow{\mathbf{h}} y} \overleftarrow{\mathbf{h}}_t + b_y$$



# Bidirectional RNN

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units:  $\overrightarrow{\mathbf{h}}$  and  $\overleftarrow{\mathbf{h}}$

outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

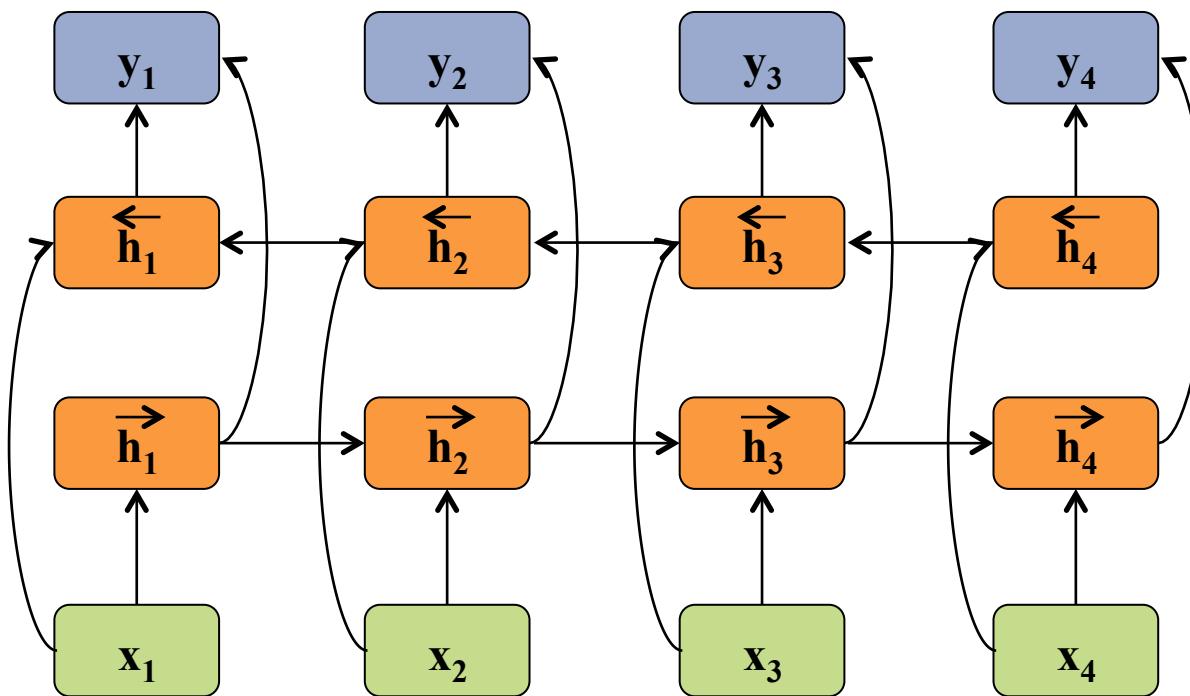
nonlinearity:  $\mathcal{H}$

Recursive Definition:

$$\overrightarrow{\mathbf{h}}_t = \mathcal{H} \left( W_{x \overrightarrow{\mathbf{h}}} x_t + W_{\overrightarrow{\mathbf{h}} \overrightarrow{\mathbf{h}}} \overrightarrow{\mathbf{h}}_{t-1} + b_{\overrightarrow{\mathbf{h}}} \right)$$

$$\overleftarrow{\mathbf{h}}_t = \mathcal{H} \left( W_{x \overleftarrow{\mathbf{h}}} x_t + W_{\overleftarrow{\mathbf{h}} \overleftarrow{\mathbf{h}}} \overleftarrow{\mathbf{h}}_{t+1} + b_{\overleftarrow{\mathbf{h}}} \right)$$

$$y_t = W_{\overrightarrow{\mathbf{h}} y} \overrightarrow{\mathbf{h}}_t + W_{\overleftarrow{\mathbf{h}} y} \overleftarrow{\mathbf{h}}_t + b_y$$



# Bidirectional RNN

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units:  $\overrightarrow{\mathbf{h}}$  and  $\overleftarrow{\mathbf{h}}$

outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

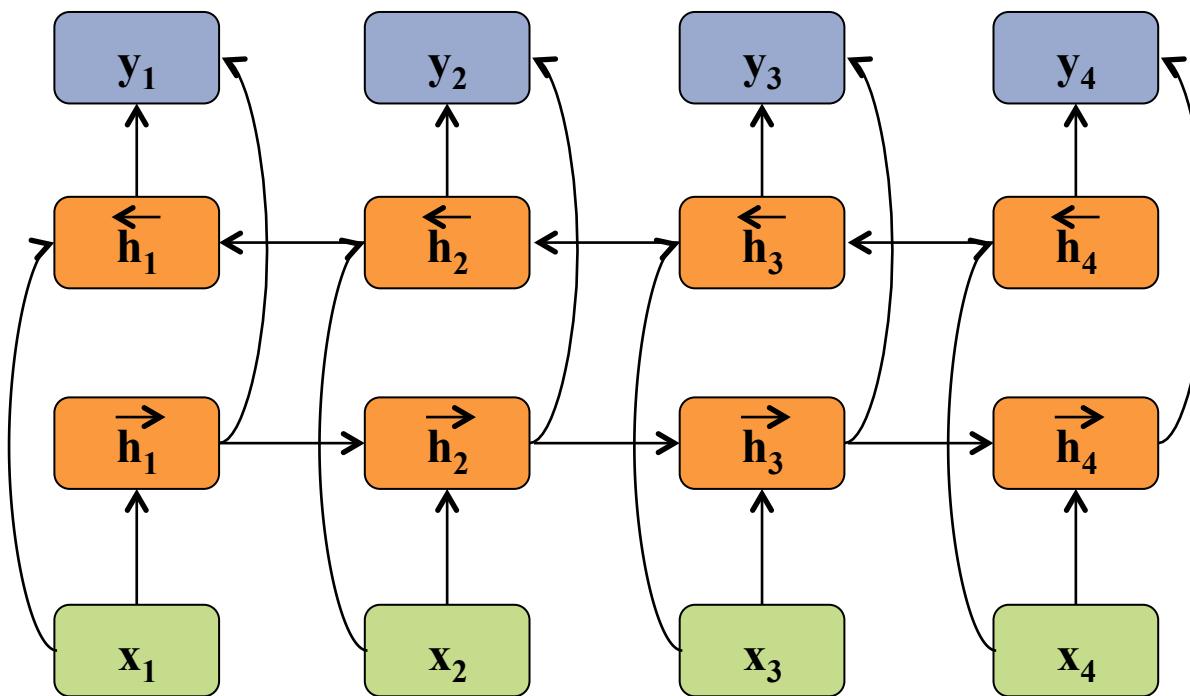
nonlinearity:  $\mathcal{H}$

Recursive Definition:

$$\overrightarrow{\mathbf{h}}_t = \mathcal{H} \left( W_{x \overrightarrow{\mathbf{h}}} x_t + W_{\overrightarrow{\mathbf{h}} \overrightarrow{\mathbf{h}}} \overrightarrow{\mathbf{h}}_{t-1} + b_{\overrightarrow{\mathbf{h}}} \right)$$

$$\overleftarrow{\mathbf{h}}_t = \mathcal{H} \left( W_{x \overleftarrow{\mathbf{h}}} x_t + W_{\overleftarrow{\mathbf{h}} \overleftarrow{\mathbf{h}}} \overleftarrow{\mathbf{h}}_{t+1} + b_{\overleftarrow{\mathbf{h}}} \right)$$

$$y_t = W_{\overrightarrow{\mathbf{h}} y} \overrightarrow{\mathbf{h}}_t + W_{\overleftarrow{\mathbf{h}} y} \overleftarrow{\mathbf{h}}_t + b_y$$



# Bidirectional RNN

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

hidden units:  $\overrightarrow{\mathbf{h}}$  and  $\overleftarrow{\mathbf{h}}$

outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

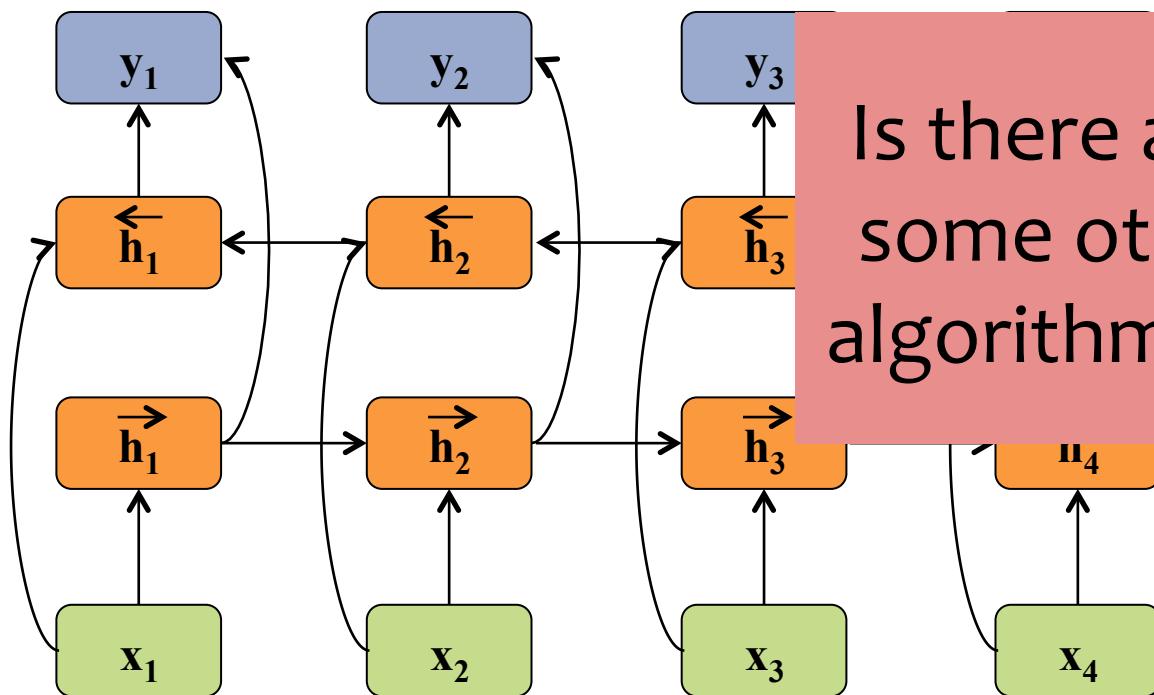
nonlinearity:  $\mathcal{H}$

Recursive Definition:

$$\overrightarrow{\mathbf{h}}_t = \mathcal{H} \left( W_{x \rightarrow h} x_t + W_{\overrightarrow{h} \rightarrow h} \overrightarrow{\mathbf{h}}_{t-1} + b_{\rightarrow h} \right)$$

$$\overleftarrow{\mathbf{h}}_t = \mathcal{H} \left( W_{x \overleftarrow{h}} x_t + W_{\overleftarrow{h} \overleftarrow{h}} \overleftarrow{\mathbf{h}}_{t+1} + b_{\overleftarrow{h}} \right)$$

$$y_t = W_{\overrightarrow{h} y} \overrightarrow{\mathbf{h}}_t + W_{\overleftarrow{h} y} \overleftarrow{\mathbf{h}}_t + b_y$$



Is there an analogy to some other recursive algorithm(s) we know?

# Deep RNNs

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

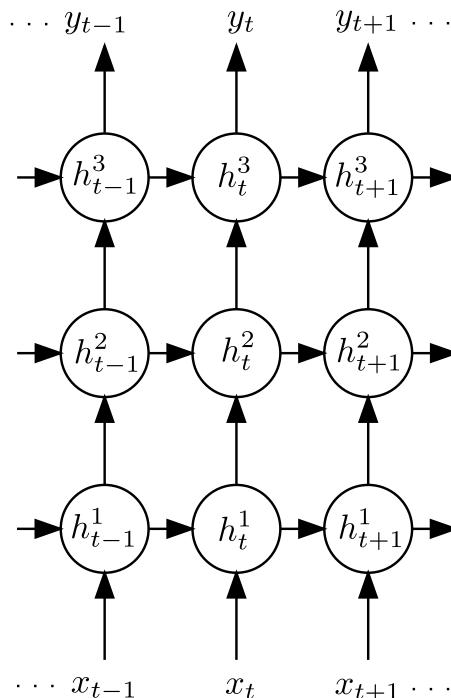
outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

nonlinearity:  $\mathcal{H}$

Recursive Definition:

$$h_t^n = \mathcal{H}(W_{h^{n-1}h^n}h_t^{n-1} + W_{h^n h^n}h_{t-1}^n + b_h^n)$$

$$y_t = W_{h^N y}h_t^N + b_y$$



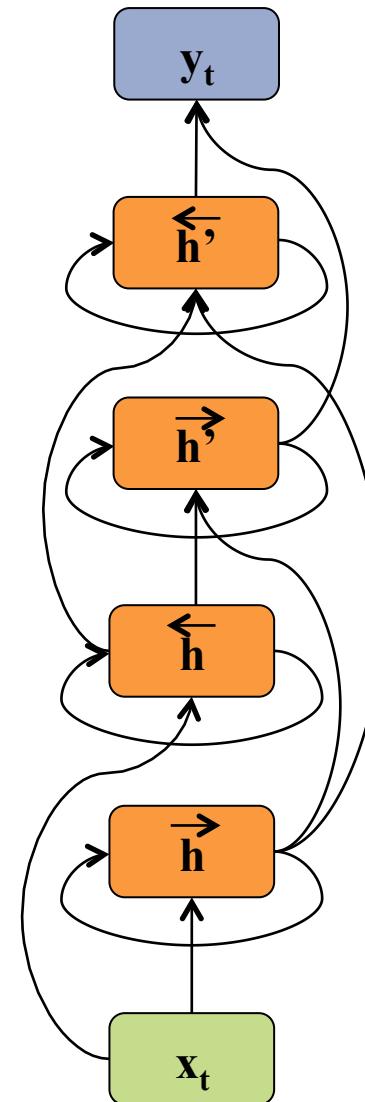
# Deep Bidirectional RNNs

inputs:  $\mathbf{x} = (x_1, x_2, \dots, x_T), x_i \in \mathcal{R}^I$

outputs:  $\mathbf{y} = (y_1, y_2, \dots, y_T), y_i \in \mathcal{R}^K$

nonlinearity:  $\mathcal{H}$

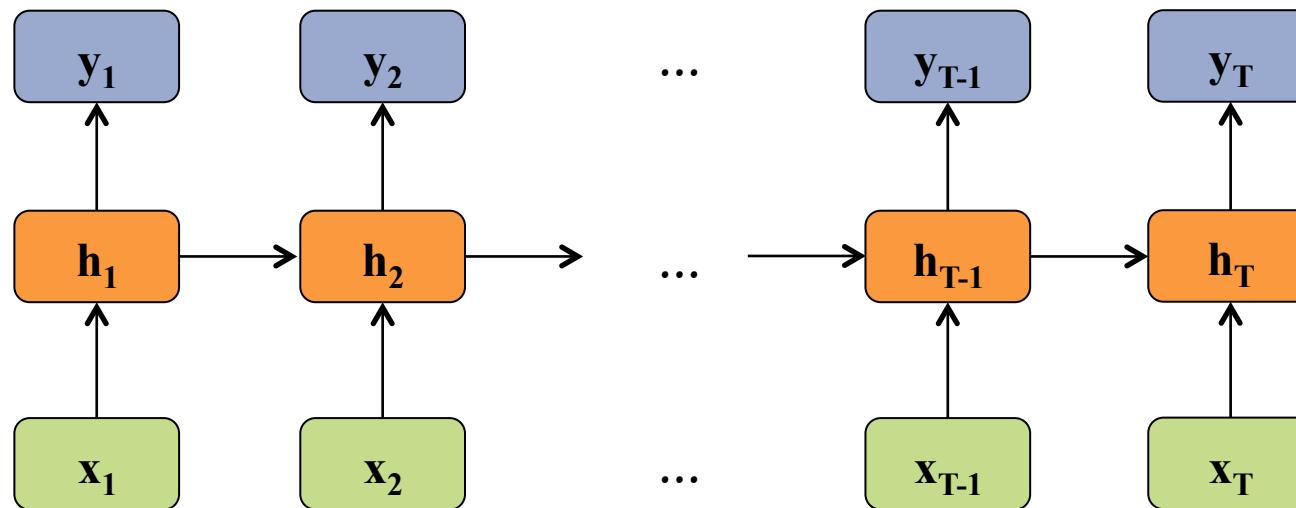
- Notice that the upper level hidden units have input from **two previous layers** (i.e. wider input)
- Likewise for the output layer
- What analogy can we draw to DNNs, DBNs, DBMs?



# Long Short-Term Memory (LSTM)

## Motivation:

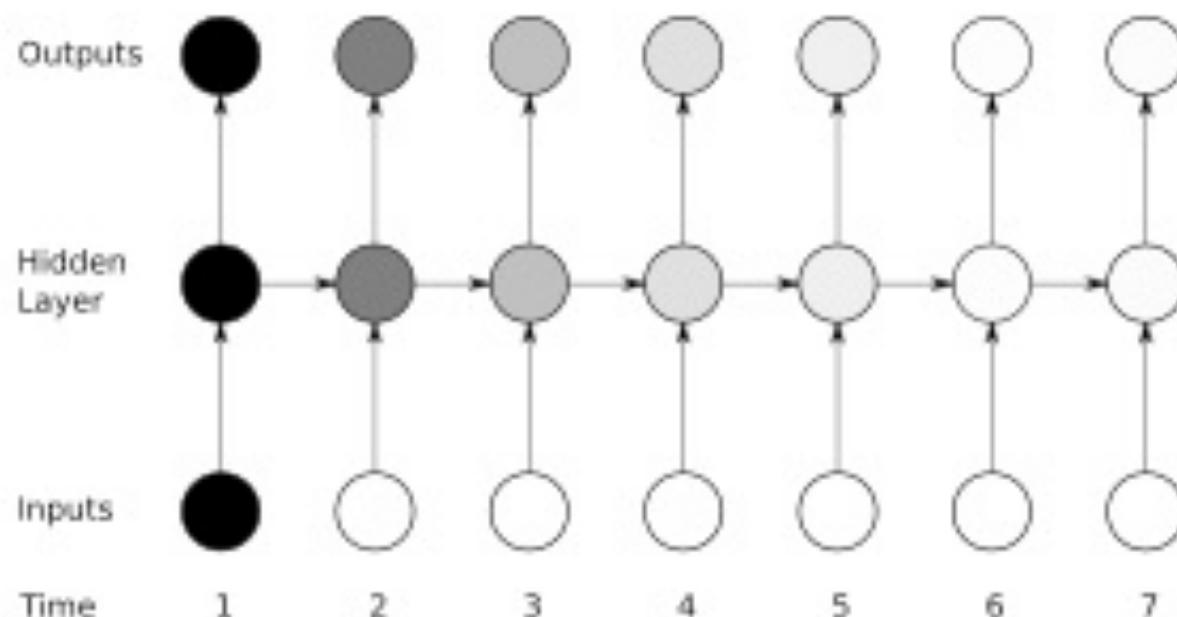
- Standard RNNs have trouble learning long distance dependencies
- LSTMs combat this issue



# Long Short-Term Memory (LSTM)

Motivation:

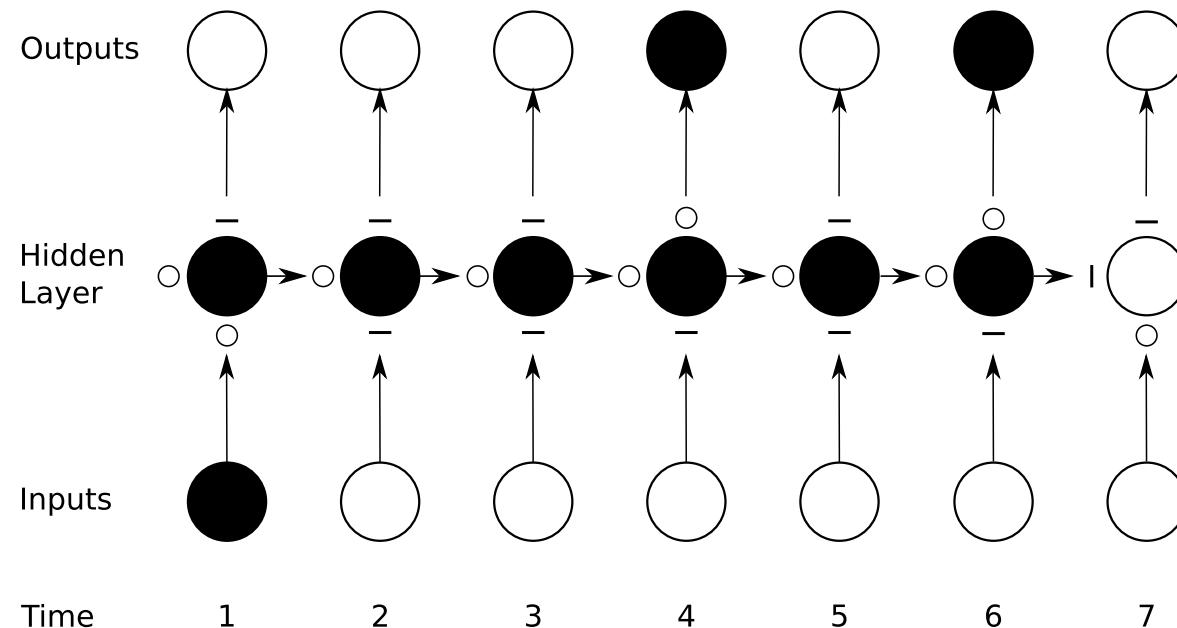
- Vanishing gradient problem for Standard RNNs
- Figure shows sensitivity (darker = more sensitive) to the input at time t=1



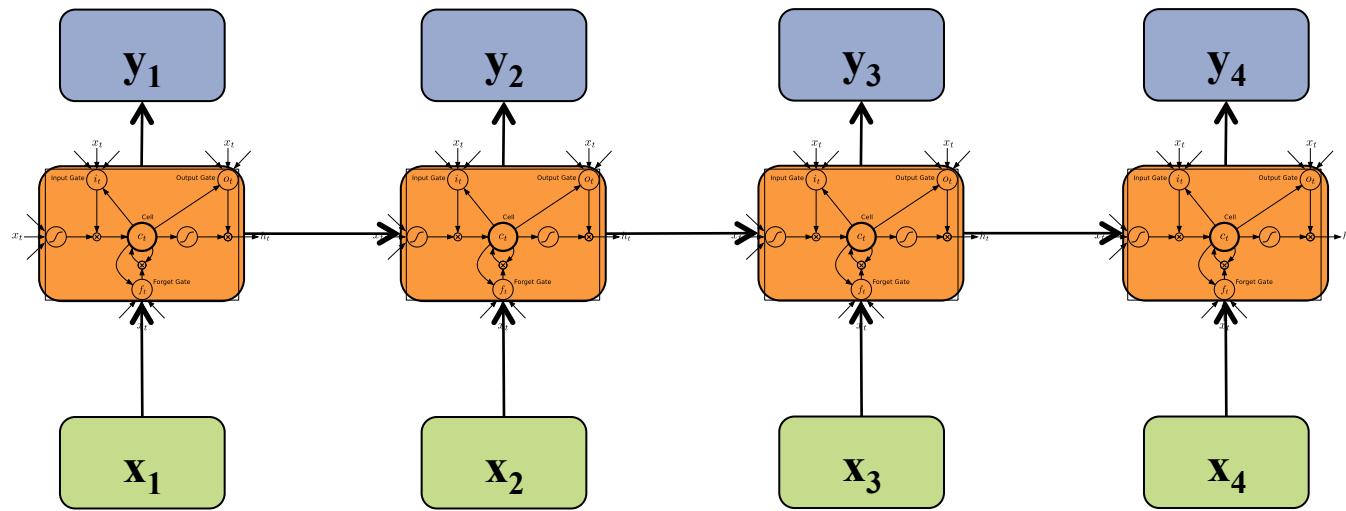
# Long Short-Term Memory (LSTM)

Motivation:

- LSTM units have a rich internal structure
- The various “gates” determine the propagation of information and can choose to “remember” or “forget” information

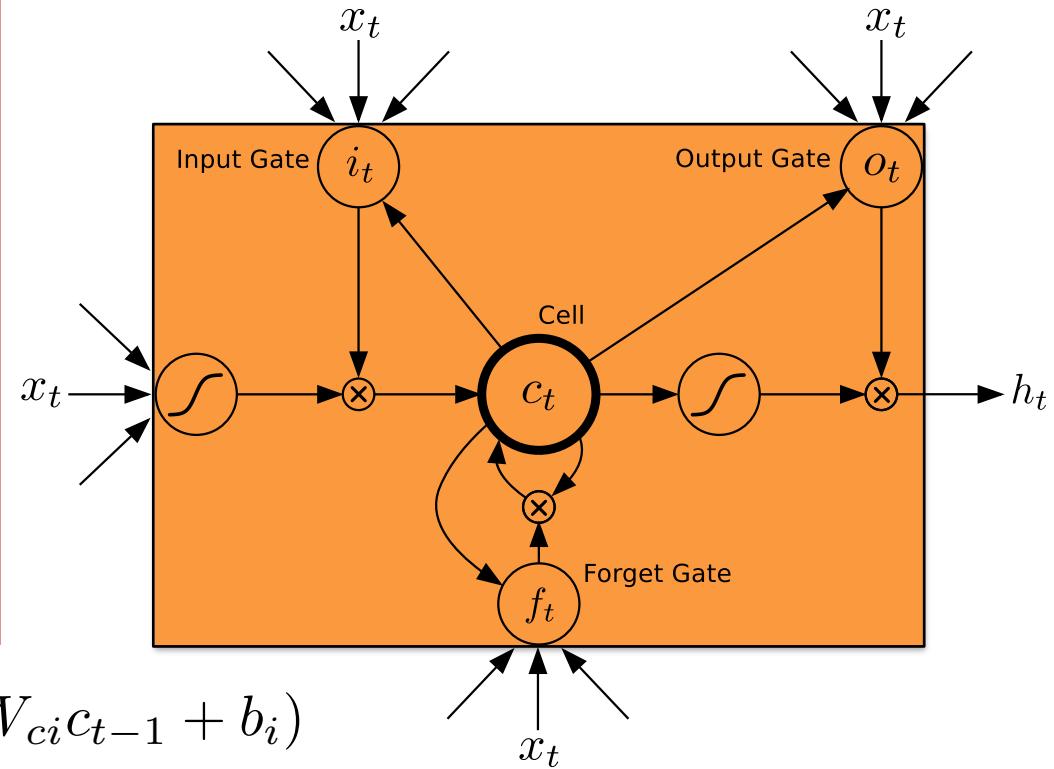


# Long Short-Term Memory (LSTM)



# Long Short-Term Memory (LSTM)

- **Input gate:** masks out the standard RNN inputs
- **Forget gate:** masks out the previous cell
- **Cell:** stores the input/forget mixture
- **Output gate:** masks out the values of the next hidden



$$i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1} + W_{ci}c_{t-1} + b_i)$$

$$f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1} + W_{cf}c_{t-1} + b_f)$$

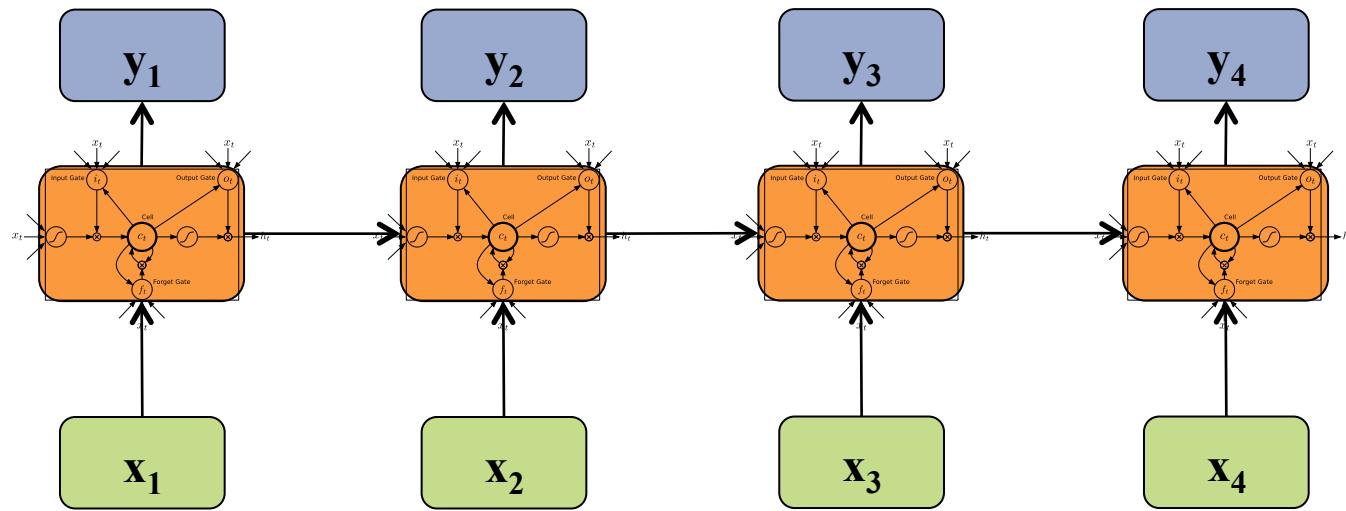
$$c_t = f_t c_{t-1} + i_t \tanh(W_{xc}x_t + W_{hc}h_{t-1} + b_c)$$

$$o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1} + W_{co}c_t + b_o)$$

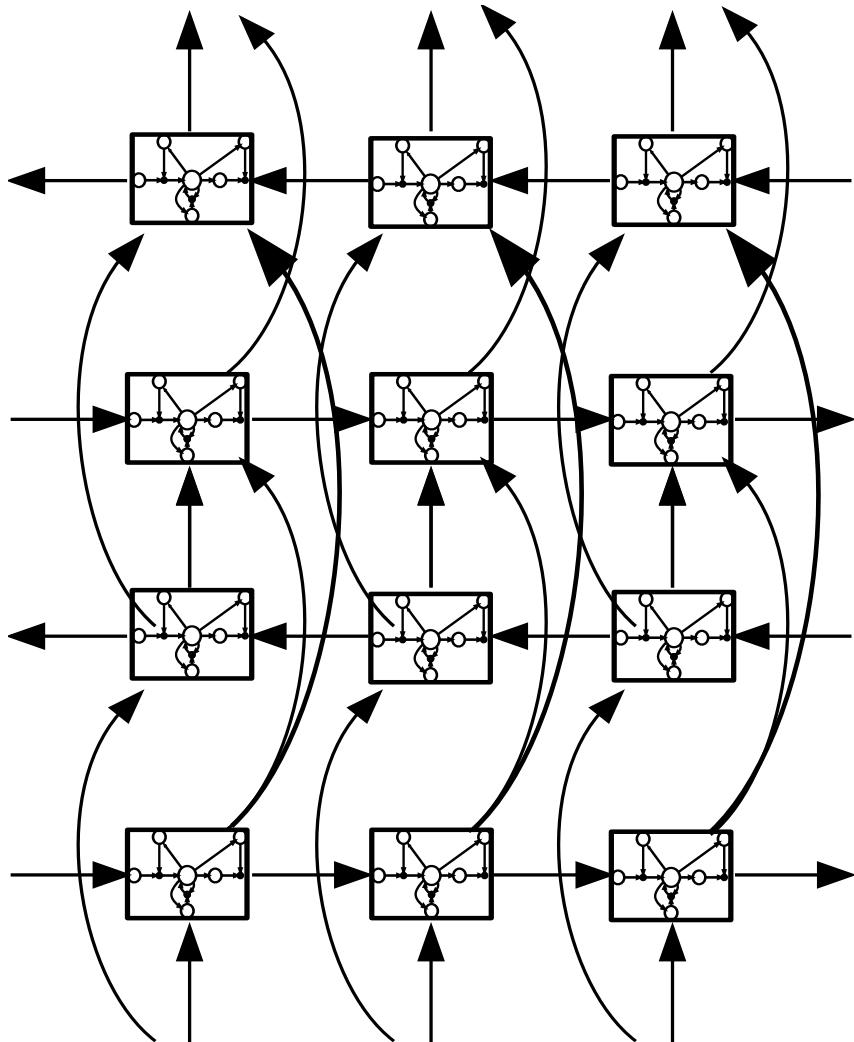
$$h_t = o_t \tanh(c_t)$$

Figure from (Graves et al., 2013)

# Long Short-Term Memory (LSTM)

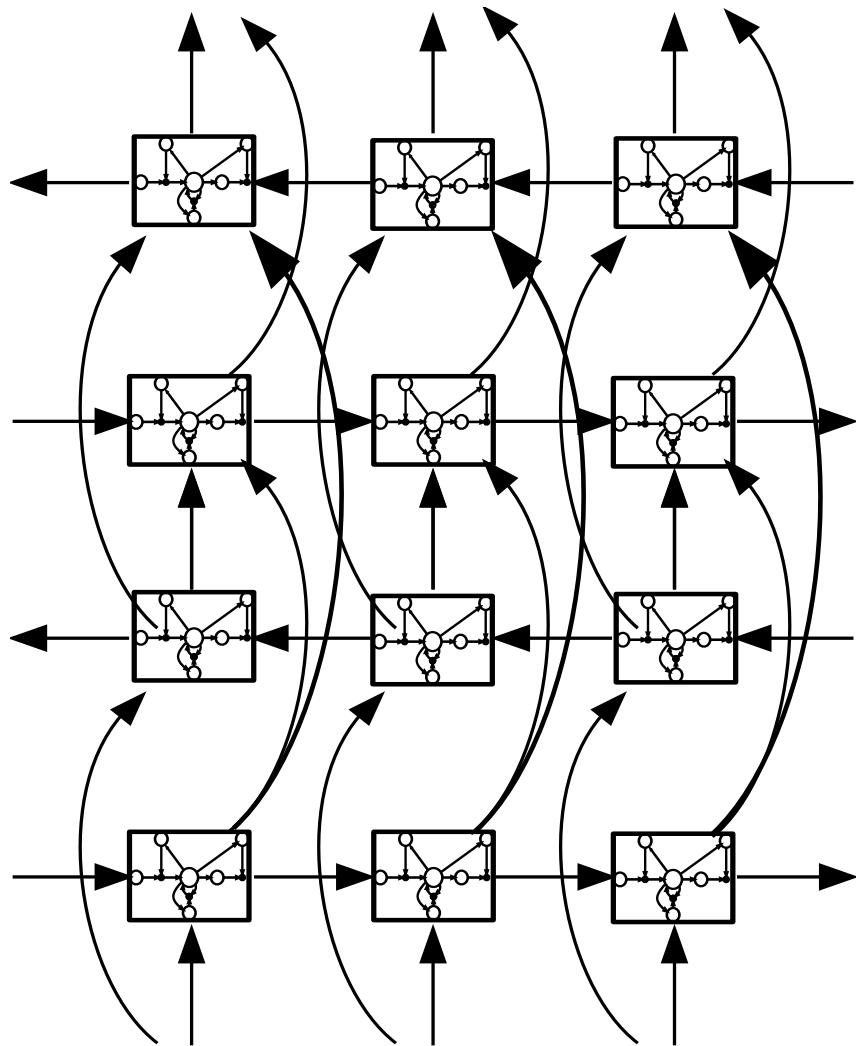


# Deep Bidirectional LSTM (DBLSTM)



- Figure: input/output layers not shown
- **Same general topology** as a Deep Bidirectional RNN, but with **LSTM units** in the hidden layers
- No additional **representational power** over DBRNN, but **easier to learn** in practice

# Deep Bidirectional LSTM (DBLSTM)



How important is this particular architecture?

Jozefowicz et al. (2015) evaluated 10,000 different LSTM-like architectures and found several variants that worked just as well on several tasks.

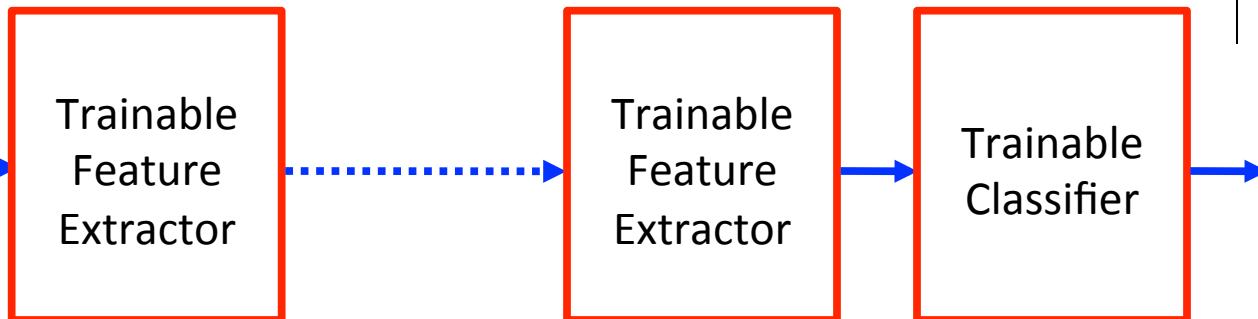
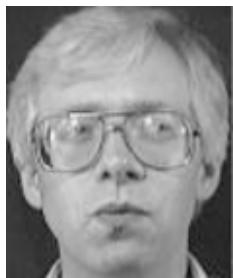
# **CONVOLUTIONAL NEURAL NETS**



# Expressive Capabilities of ANNs

- Boolean functions:
  - Every Boolean function can be represented by network with single hidden layer
  - But might require exponential (in number of inputs) hidden units
- Continuous functions:
  - Every bounded continuous function can be approximated with arbitrary small error, by network with one hidden layer [Cybenko 1989; Hornik et al 1989]
  - Any function can be approximated to arbitrary accuracy by a network with two hidden layers [Cybenko 1988].

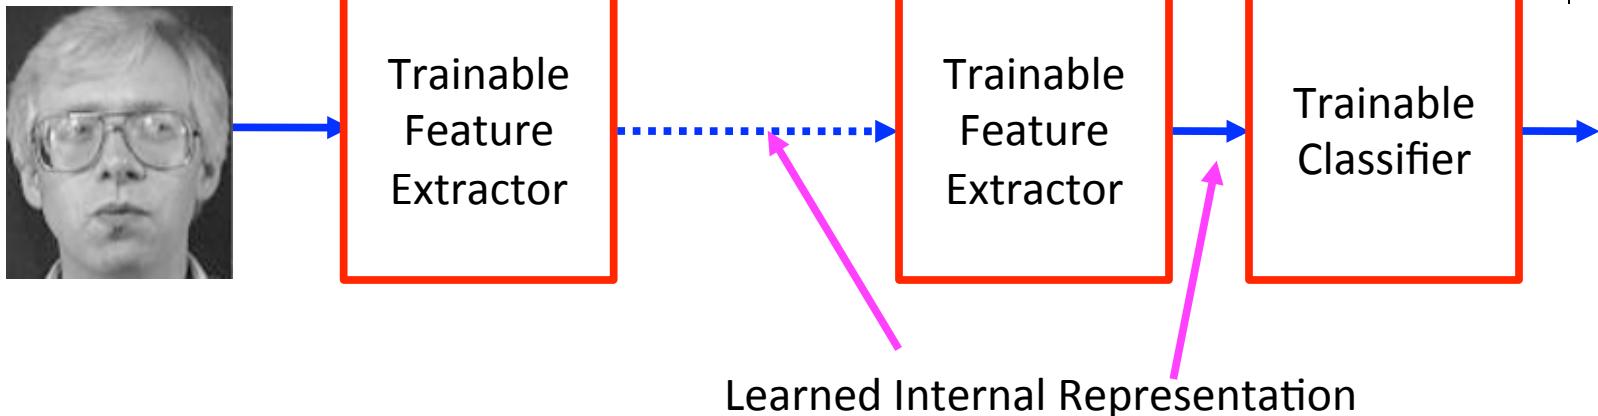
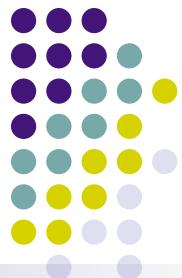
# Using ANN to hierarchical representation



**Good Representations are hierarchical**

- In Language: hierarchy in syntax and semantics
  - Words->Parts of Speech->Sentences->Text
  - Objects,Actions,Attributes...-> Phrases -> Statements -> Stories
- In Vision: part-whole hierarchy
  - Pixels->Edges->Textons->Parts->Objects->Scenes

# “Deep” learning: learning hierarchical representations



- Deep Learning: learning a hierarchy of internal representations
- From low-level features to mid-level invariant representations, to object identities
- Representations are increasingly invariant as we go up the layers
- using multiple stages gets around the specificity/invariance dilemma

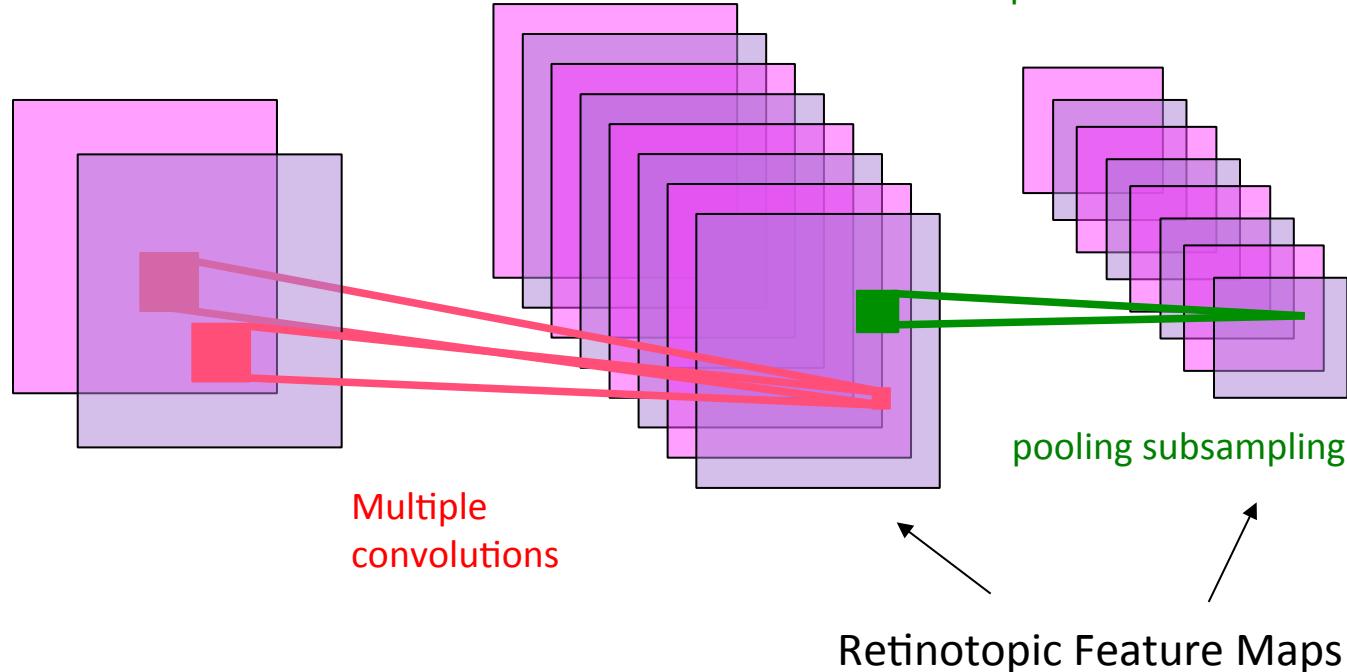


# Filtering+NonLinearity+Pooling = 1 stage of a Convolutional Net

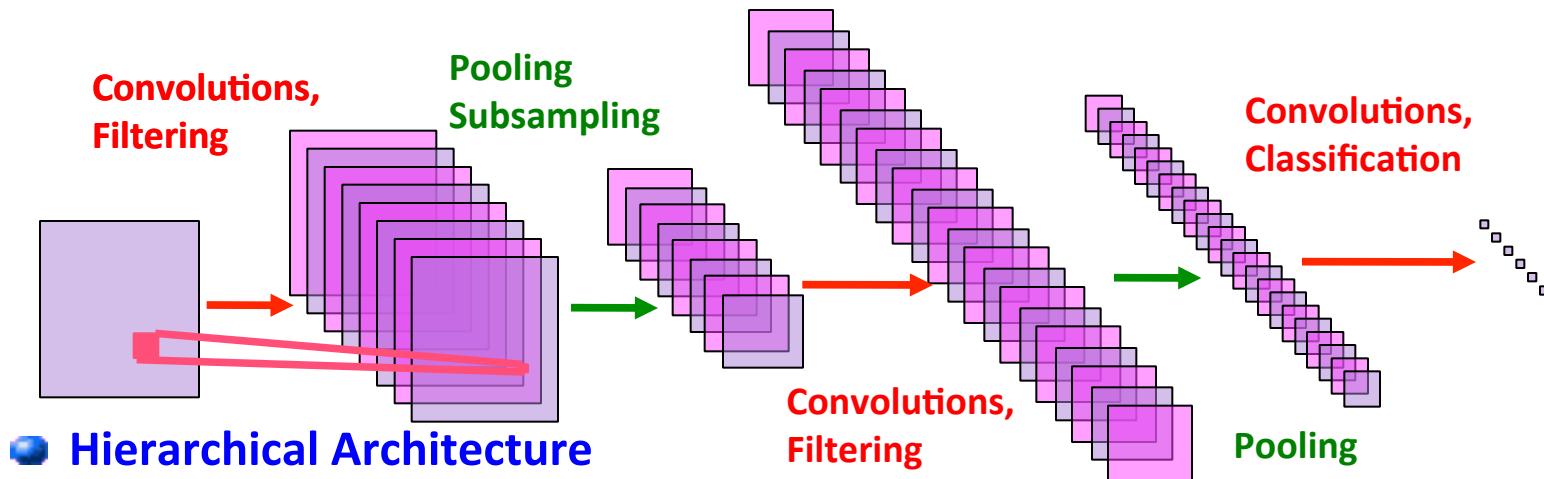
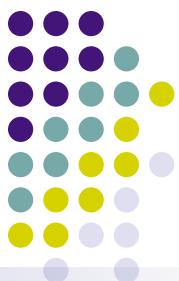
- [Hubel & Wiesel 1962]:
  - simple cells detect local features
  - complex cells “pool” the outputs of simple cells within a retinotopic neighborhood.

“Simple cells”

“Complex cells”



# Convolutional Network: Multi-Stage Trainable Architecture



- **Hierarchical Architecture**

- ▶ Representations are more global, more invariant, and more abstract as we go up the layers

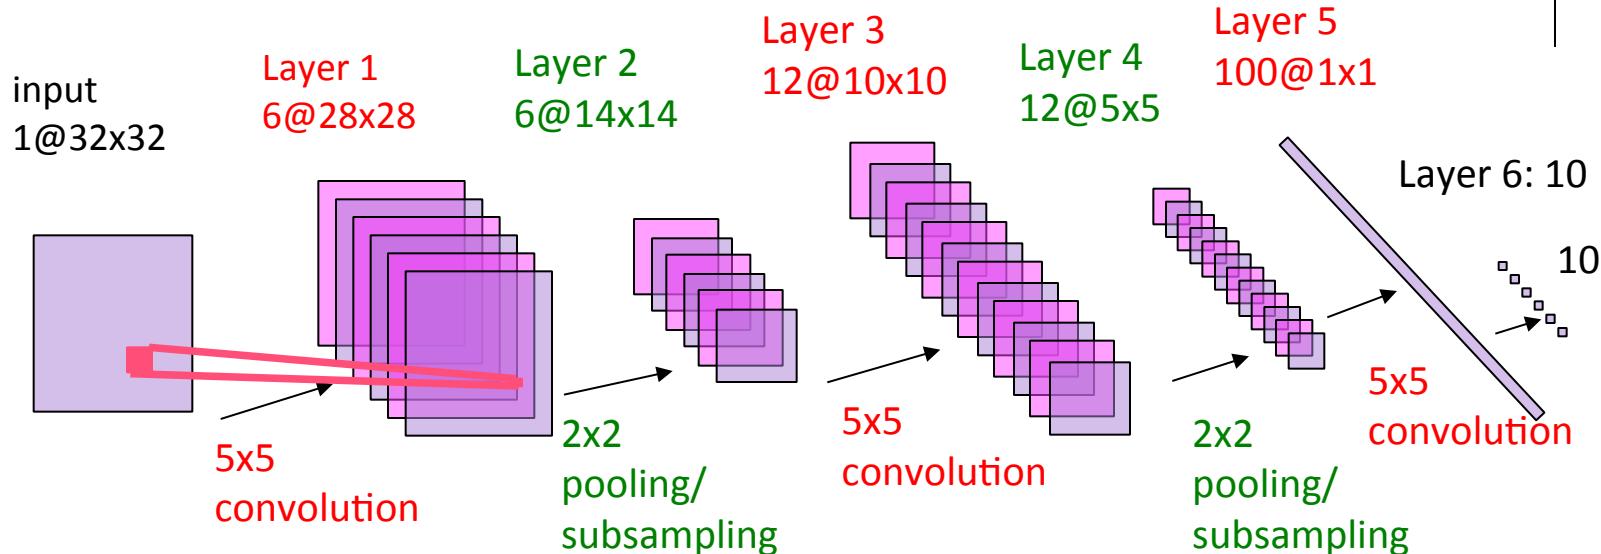
- **Alternated Layers of Filtering and Spatial Pooling**

- ▶ Filtering detects conjunctions of features
- ▶ Pooling computes local disjunctions of features

- **Fully Trainable**

- ▶ All the layers are trainable

# Convolutional Net Architecture for Hand-writing recognition

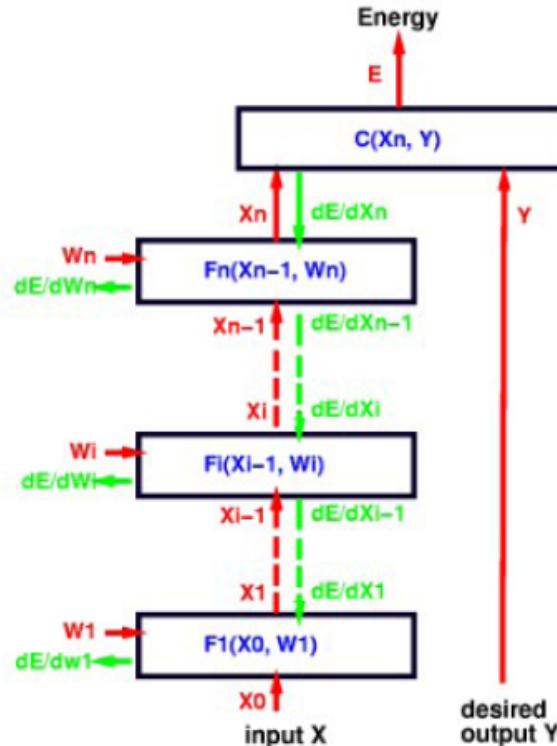


- Convolutional net for handwriting recognition (400,000 synapses)
  - Convolutional layers (simple cells): all units in a feature plane share the same weights
  - Pooling/subsampling layers (complex cells): for invariance to small distortions.
  - Supervised gradient-descent learning using back-propagation
  - The entire network is trained end-to-end. All the layers are trained simultaneously.
  - [LeCun et al. Proc IEEE, 1998]



# How to train?

To compute all the derivatives, we use a backward sweep called the **back-propagation algorithm** that uses the recurrence equation for  $\frac{\partial E}{\partial X_i}$

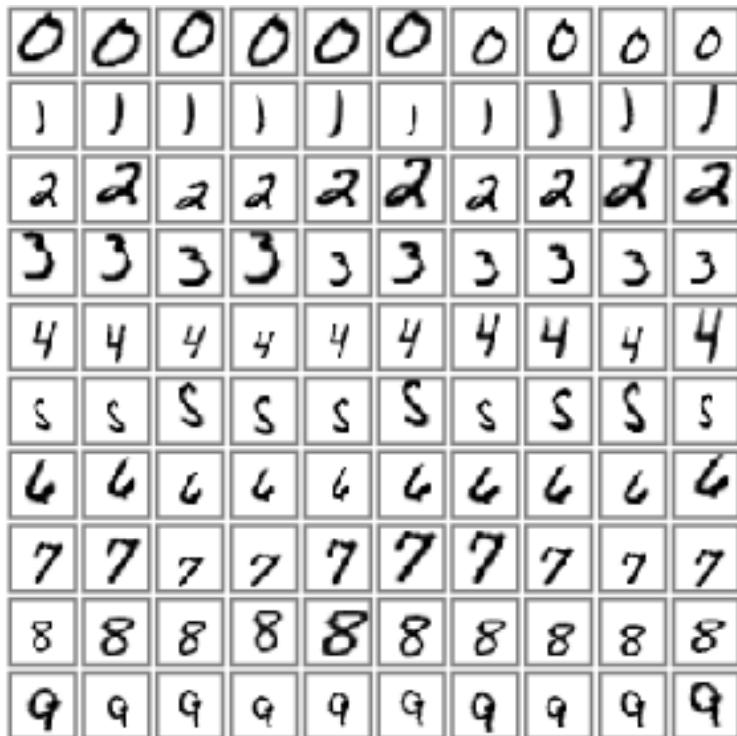


- $\frac{\partial E}{\partial X_n} = \frac{\partial C(X_n, Y)}{\partial X_n}$
- $\frac{\partial E}{\partial X_{n-1}} = \frac{\partial E}{\partial X_n} \frac{\partial F_n(X_{n-1}, W_n)}{\partial X_{n-1}}$
- $\frac{\partial E}{\partial W_n} = \frac{\partial E}{\partial X_n} \frac{\partial F_n(X_{n-1}, W_n)}{\partial W_n}$
- $\frac{\partial E}{\partial X_{n-2}} = \frac{\partial E}{\partial X_{n-1}} \frac{\partial F_{n-1}(X_{n-2}, W_{n-1})}{\partial X_{n-2}}$
- $\frac{\partial E}{\partial W_{n-1}} = \frac{\partial E}{\partial X_{n-1}} \frac{\partial F_{n-1}(X_{n-2}, W_{n-1})}{\partial W_{n-1}}$
- ....etc, until we reach the first module.
- we now have all the  $\frac{\partial E}{\partial W_i}$  for  $i \in [1, n]$ .



# Application: MNIST Handwritten Digit Dataset

3 6 8 1 7 9 6 6 4 1  
6 7 5 7 8 6 3 4 8 5  
2 1 7 9 7 1 2 8 4 6  
4 8 1 9 0 1 8 8 9 4  
7 6 1 8 6 4 1 5 6 0  
7 5 9 2 6 5 8 1 9 7  
2 2 2 2 2 3 4 4 8 0  
0 2 3 8 0 7 3 8 5 7  
0 1 4 6 4 6 0 2 4 3  
7 1 2 8 1 6 9 8 6 1



Handwritten Digit Dataset MNIST: 60,000 training samples, 10,000 test samples

# Results on MNIST Handwritten Digits

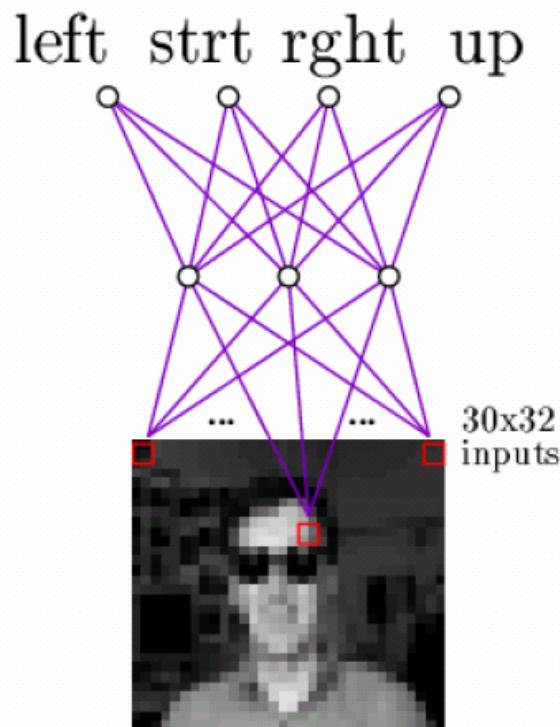


CLASSIFIER	DEFORMATION	PREPROCESSING	ERROR (%)	Reference
linear classifier (1-layer NN)		none	12.00	LeCun et al. 1998
linear classifier (1-layer NN)		deskewing	8.40	LeCun et al. 1998
pairwise linear classifier		deskewing	7.60	LeCun et al. 1998
K-nearest-neighbors, (L2)		none	3.09	Kenneth Wilder, U. Chicago
K-nearest-neighbors, (L2)		deskewing	2.40	LeCun et al. 1998
K-nearest-neighbors, (L2)		deskew, clean, blur	1.80	Kenneth Wilder, U. Chicago
K-NN L3, 2 pixel jitter		deskew, clean, blur	1.22	Kenneth Wilder, U. Chicago
<b>K-NN, shape context matching</b>		<b>shape context feature</b>	<b>0.63</b>	<b>Belongie et al. IEEE PAMI 2002</b>
40 PCA + quadratic classifier		none	3.30	LeCun et al. 1998
1000 RBF + linear classifier		none	3.60	LeCun et al. 1998
K-NN, Tangent Distance		subsample 16x16 pixels	1.10	LeCun et al. 1998
SVM, Gaussian Kernel		none	1.40	
SVM deg 4 polynomial		deskewing	1.10	LeCun et al. 1998
Reduced Set SVM deg 5 poly		deskewing	1.00	LeCun et al. 1998
Virtual SVM deg-9 poly	Affine	none	0.80	LeCun et al. 1998
V-SVM, 2-pixel jittered		none	0.68	DeCoste and Scholkopf, MLJ 2002
<b>V-SVM, 2-pixel jittered</b>		<b>deskewing</b>	<b>0.56</b>	<b>DeCoste and Scholkopf, MLJ 2002</b>
2-layer NN, 300 HU, MS E		none	4.70	LeCun et al. 1998
2-layer NN, 300 HU, MS E,	Affine	none	3.60	LeCun et al. 1998
2-layer NN, 300 HU		deskewing	1.60	LeCun et al. 1998
3-layer NN, 500+ 150 HU		none	2.95	LeCun et al. 1998
3-layer NN, 500+ 150 HU	Affine	none	2.45	LeCun et al. 1998
3-layer NN, 500+ 300 HU, CE, reg		none	1.53	Hinton, unpublished, 2005
2-layer NN, 800 HU, CE		none	1.60	Simard et al., ICDAR 2003
2-layer NN, 800 HU, CE	Affine	none	1.10	Simard et al., ICDAR 2003
2-layer NN, 800 HU, MS E	Elastic	none	0.90	Simard et al., ICDAR 2003
<b>2-layer NN, 800 HU, CE</b>	<b>Elastic</b>	<b>none</b>	<b>0.70</b>	<b>Simard et al., ICDAR 2003</b>
Convolutional net LeNet-1		subsample 16x16 pixels	1.70	LeCun et al. 1998
Convolutional net LeNet-4		none	1.10	LeCun et al. 1998
Convolutional net LeNet-5,		none	0.95	LeCun et al. 1998
<b>Conv. net LeNet-5,</b>	<b>Affine</b>	<b>none</b>	<b>0.80</b>	<b>LeCun et al. 1998</b>
Boosted LeNet-4	Affine	none	0.70	LeCun et al. 1998
<b>Conv. net, CE</b>	<b>Affine</b>	<b>none</b>	<b>0.60</b>	<b>Simard et al., ICDAR 2003</b>
<b>Conv net, CE</b>	<b>Elastic</b>	<b>none</b>	<b>0.40</b>	<b>Simard et al., ICDAR 2003</b>

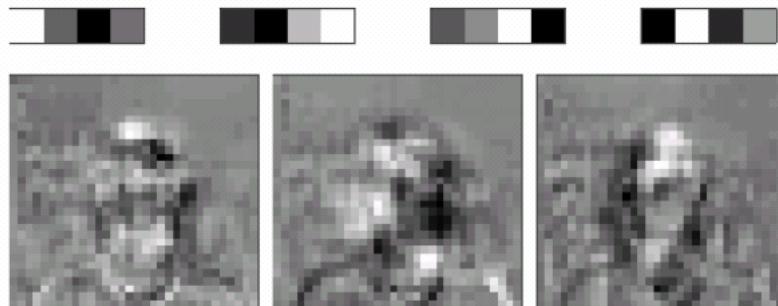


# Application: ANN for Face Reco.

- The model



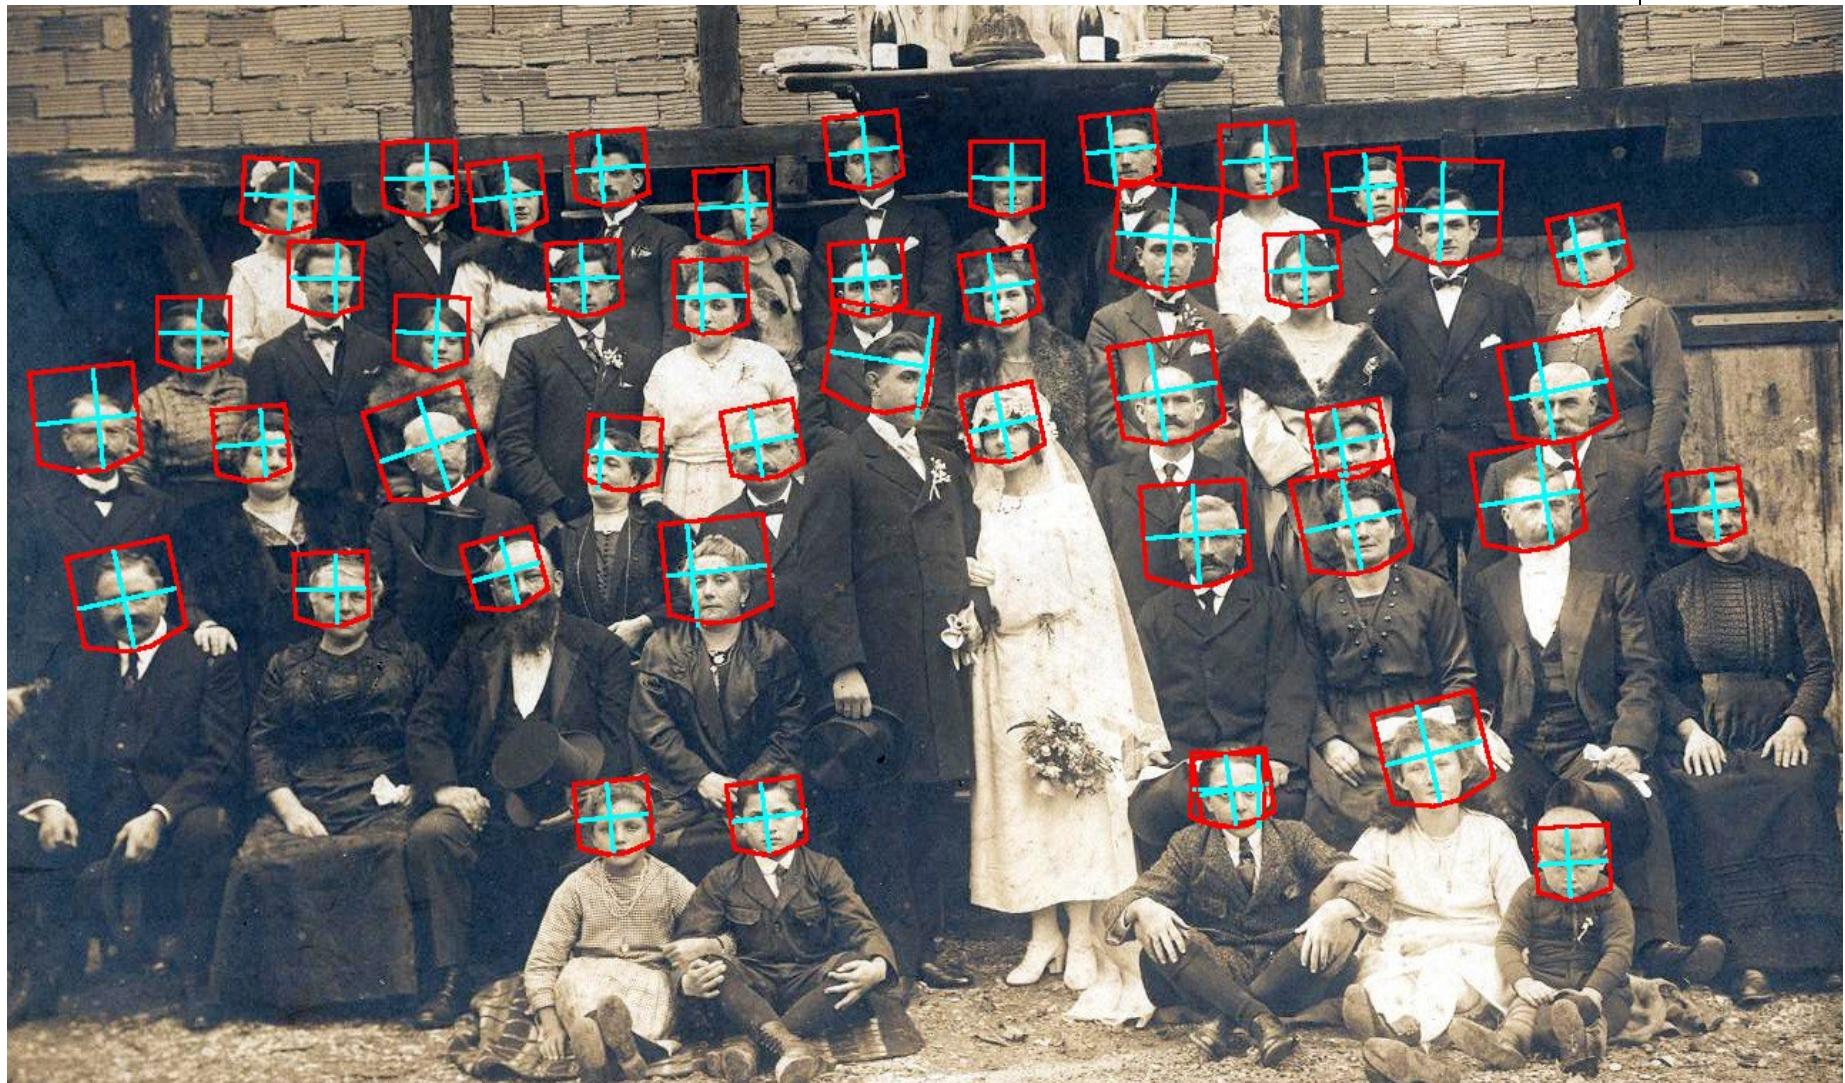
- The learned hidden unit weights



Typical input images

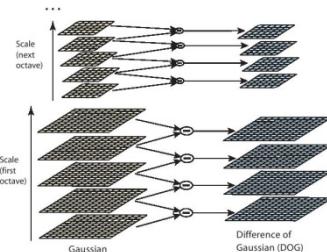
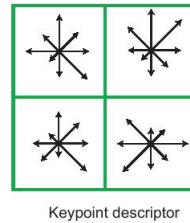
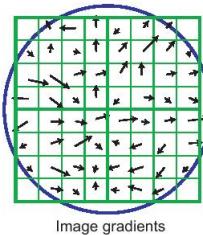
<http://www.cs.cmu.edu/~tom/faces.html>

# Face Detection with a Convolutional Net

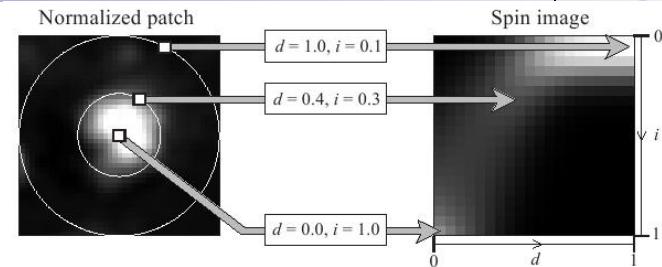
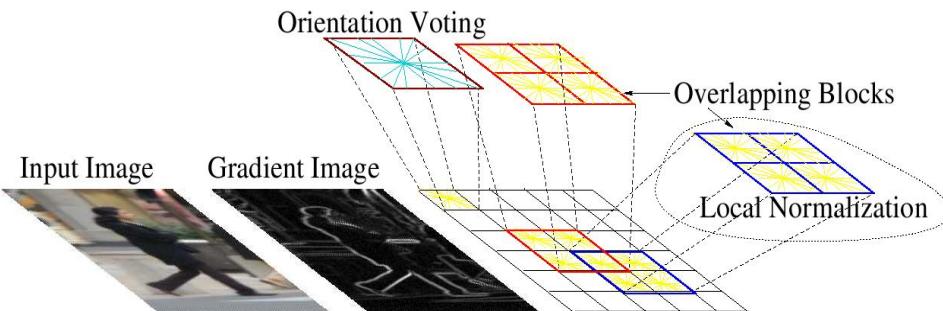




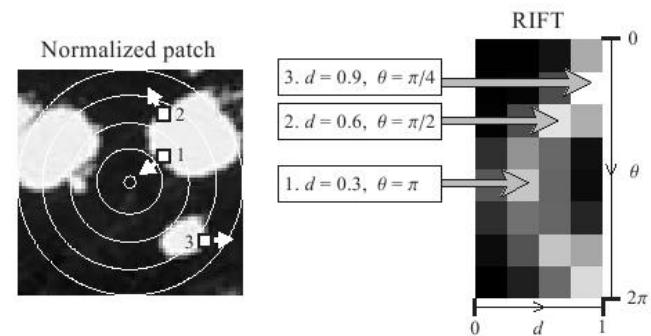
# Computer vision features



SIFT

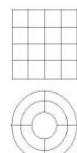


Spin image



## Drawbacks of feature engineering

1. Needs expert knowledge
2. Time consuming hand-tuning

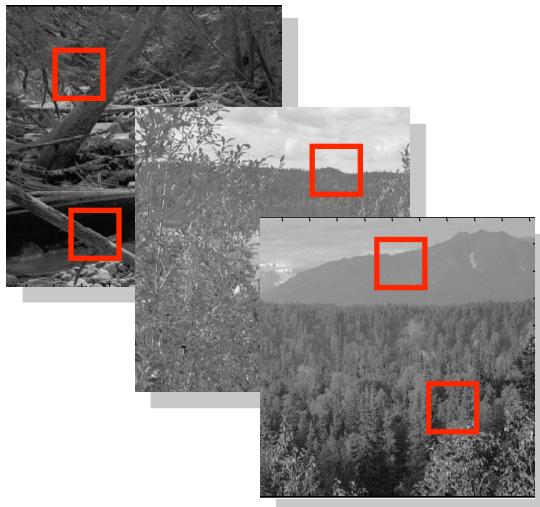


(e)

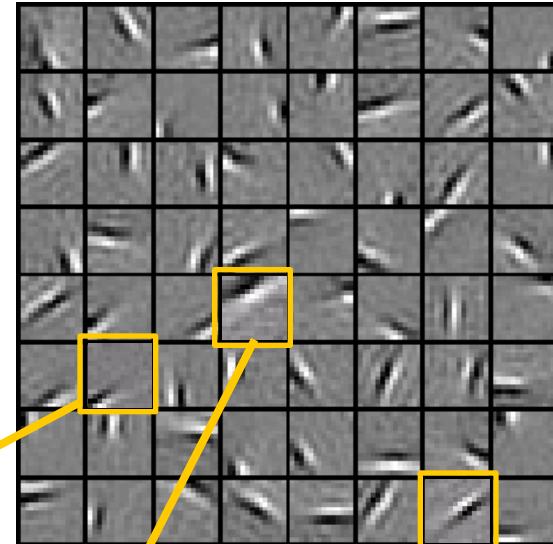


# Sparse coding on images

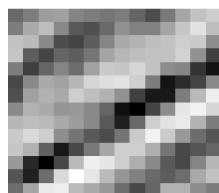
Natural Images



Learned bases: “Edges”



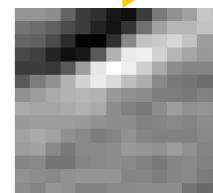
New example



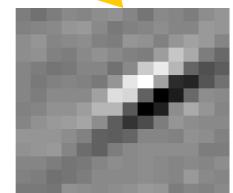
$$= 0.8 *$$



$$+ 0.3 *$$



$$+ 0.5 *$$

 $x$ 

$$= 0.8 *$$

 $b_{36}$ 

$$+ 0.3 *$$

 $b_{42}$ 

$$+ 0.5 *$$

 $b_{65}$ 

[0, 0, ... 0.8, ..., 0.3, ..., 0.5, ...] = coefficients (feature representation)

Courtesy: Lee and Ng  
75



# Basis (or features) can be learned by Optimization

Given input data  $\{x^{(1)}, \dots, x^{(m)}\}$ , we want to find good bases  $\{b_1, \dots, b_n\}$ :

$$\min_{b,a} \sum_i \left\| x^{(i)} - \underbrace{\sum_j a_j^{(i)} b_j}_{\text{Reconstruction error}} \right\|_2^2 + \beta \underbrace{\sum_i \|a^{(i)}\|_1}_{\text{Sparsity penalty}}$$

Reconstruction error

Sparsity penalty

$$\forall j: \|b_j\| \leq 1$$

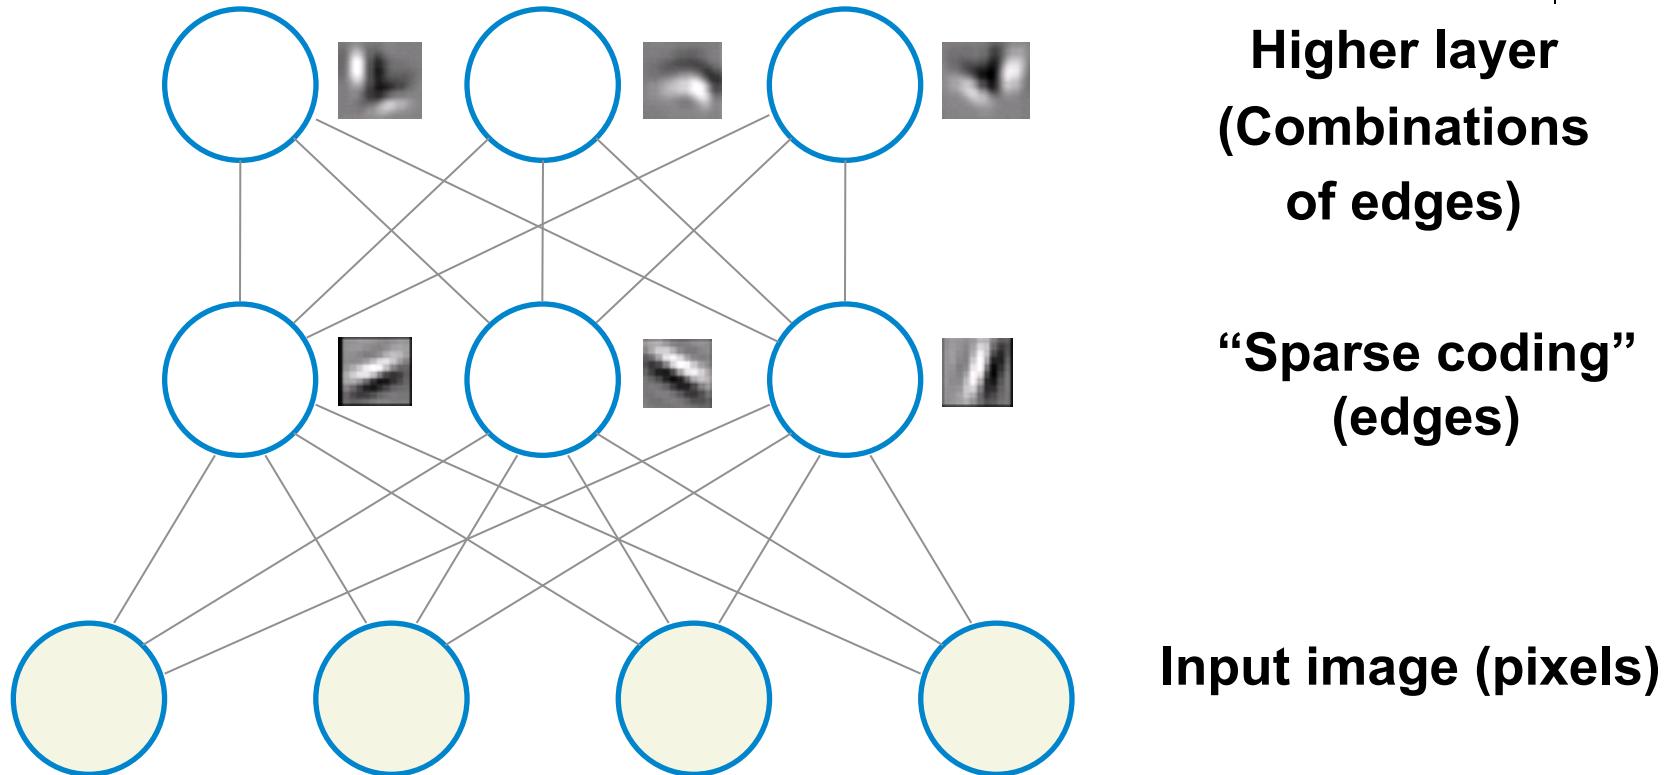
Normalization  
constraint

Solve by alternating minimization:

- Keep  $b$  fixed, find optimal  $a$ .
- Keep  $a$  fixed, find optimal  $b$ .



# Learning Feature Hierarchy

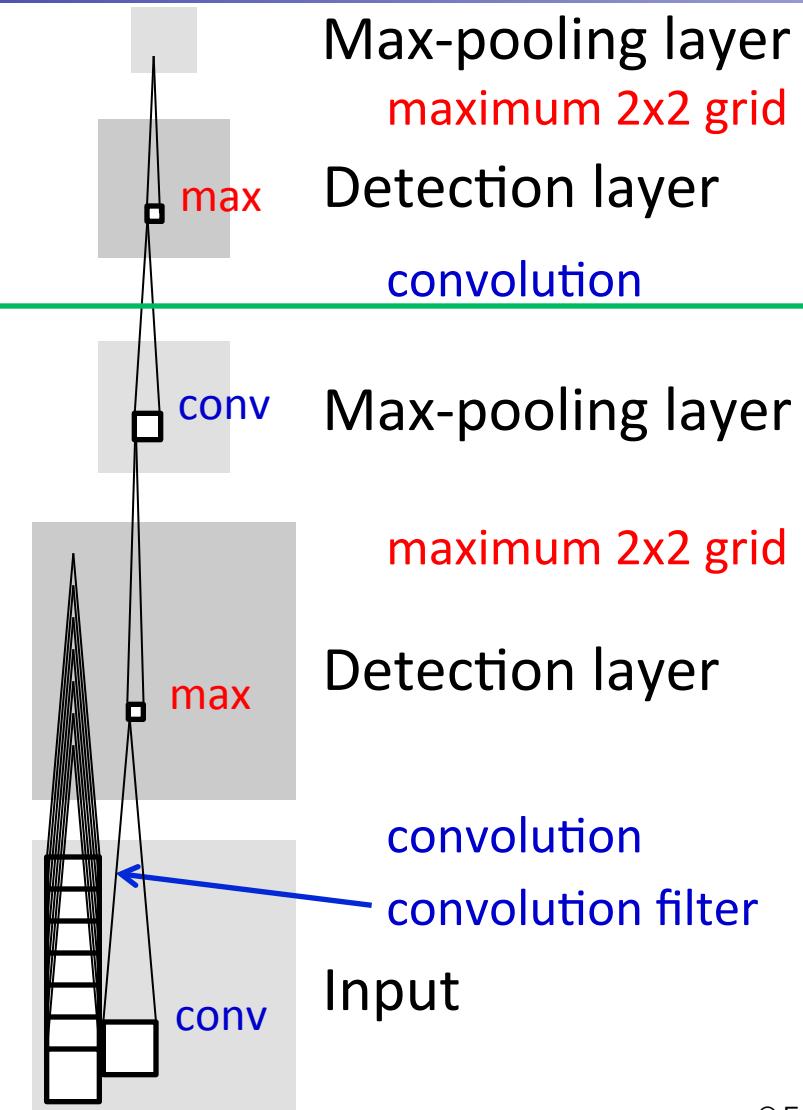


DBN (Hinton et al., 2006) with additional sparseness constraint.

[Related work: Hinton, Bengio, LeCun, and others.]



# Convolutional architectures



Max-pooling layer

maximum 2x2 grid

Detection layer

convolution

conv

Max-pooling layer

maximum 2x2 grid

max

Detection layer

convolution

convolution filter

Input

conv

- Weight sharing by convolution (e.g., [Lecun et al., 1989])
- “Max-pooling”  
Invariance  
Computational efficiency  
Deterministic and feed-forward
- One can develop convolutional Restricted Boltzmann machine (CRBM).
- One can define *probabilistic max-pooling* that *combine bottom-up and top-down information*.

# Convolutional Deep Belief Networks

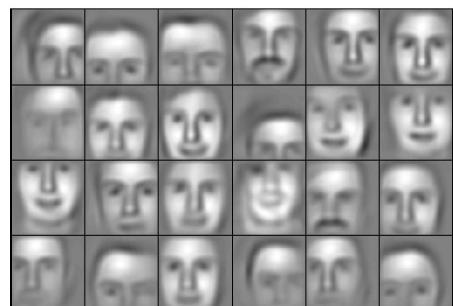


- Bottom-up (greedy), layer-wise training
  - Train one layer (convolutional RBM) at a time.
- Inference (approximate)
  - Undirected connections for all layers (Markov net)  
[Related work: Salakhutdinov and Hinton, 2009]
  - Block Gibbs sampling or mean-field
  - Hierarchical probabilistic inference

# Unsupervised learning of object-parts



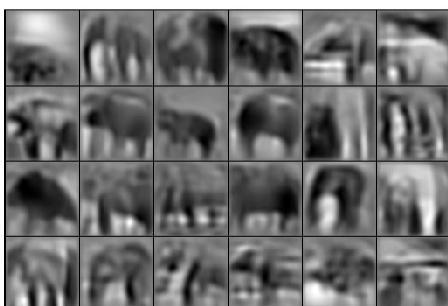
Faces



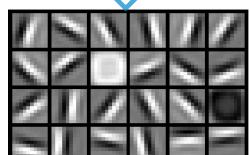
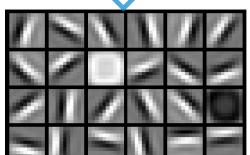
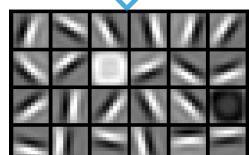
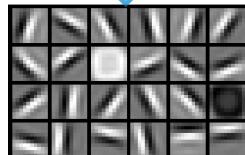
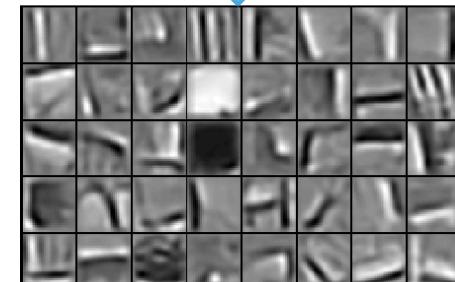
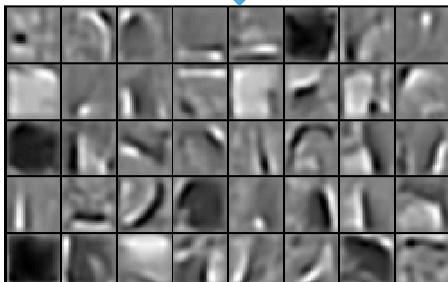
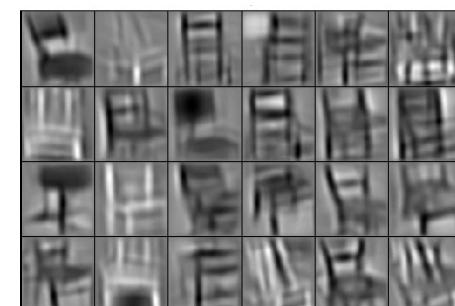
Cars



Elephants



Chairs





# Weaknesses & Criticisms

- Learning everything. Better to encode prior knowledge about structure of images.  
A: Compare with machine learning vs. linguists debate in NLP.
- Results not yet competitive with best engineered systems.  
A: Agreed. True for some domains.

# Tutorials

- LSTMs
  - Christopher Olah’s blog
  - <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- Convolutional Neural Networks
  - Andrej Karpathy, CS231n Notes
  - <http://cs231n.github.io/convolutional-networks/>